

**INTERNATIONAL ISLAMIC UNIVERSITY
CHITTAGONG**

**CacheFlow: Memory-Aware Expert Scheduling
for Efficient Mixture-of-Experts Models**

Supervised by:

Jamil As Ad

Lecturer

Department of Computer Science and Engineering
International Islamic University Chittagong

Submitted By:

Mir Md. Tarhimul Quader C221017

Md. Iftekar Ahamed Sayem C221020

Turja Dutta C221026

**BACHELOR OF SCIENCE IN COMPUTER SCIENCE
AND ENGINEERING**

Department of Computer Science and Engineering
International Islamic University Chittagong

February 2026

Declaration

We hereby affirm the following statements regarding our thesis:

1. The thesis has been successfully completed as part of our undergraduate degree program at International Islamic University Chittagong.
2. The thesis work does not contain any previously published or third-party content without proper citation.
3. The thesis work has not been previously submitted for any other degree or diploma at any other university or institution.
4. We have appropriately acknowledged all significant sources of contribution in the thesis.

Student's Full Name and Metric ID:

Mir Md. Tarhimul Quader (C221017)

Md. Iftaker Ahamed Sayem (C221020)

Turja Dutta (C221026)

Supervisor's Declaration

I formally state that I have examined this thesis and claim it to be of sufficient quality and scope to be granted the undergraduate degree of Bachelor of Science in Computer Science and Engineering.

.....

Jamil As-Ad

Lecturer

Department of Computer Science and Engineering

International Islamic University Chittagong

Dedication

This thesis report is dedicated to us, our supervisor and our family. The teamwork was satisfactory and the family's support was incredibly Amazing.

Our dedicated and most hard-working Supervisor who has been a constant support throughout these months. In this document, the contributions are acknowledged too.

Acknowledgment

First of all, we are grateful to the Almighty Allah for his kindness, which allowed us to complete our thesis in spite of several challenges.

Second, we would like to express our appreciation to our supervisor **Jamil As Ad** Sir for his unwavering support and guidance from the start of our research.

Ethical Statement

Hereby we state that None of the unethical practices were used in the completion of our thesis work. The data we used for research purposes are original. We carefully checked every citation we used here. The writers of the work accept all the liabilities for any kind of violation of the thesis rule.

Abstract

Large-scale Mixture-of-Experts (MoE) models have demonstrated remarkable capabilities across natural language processing and multimodal tasks, yet their practical deployment remains severely constrained by GPU memory bottlenecks. While MoE architectures promise efficient computation through sparse expert activation, their deployment requires simultaneous residency of all expert weight matrices in GPU VRAM—a requirement that incurs prohibitive memory costs (20–30 GB for state-of-the-art models) and renders such systems inaccessible on consumer-grade and edge hardware. Existing approaches either incur out-of-memory failures (DeepSpeed MoE on constrained hardware) or introduce severe latency penalties through synchronous disk-to-GPU staging (MoE-Infinity: 920 ms/token, $\sim$ 5.4 GB pinned host memory allocation). This thesis presents CacheFlow, a comprehensive three-tier memory hierarchy system that decouples GPU memory requirements from total expert parameter count, enabling efficient MoE inference on resource-constrained hardware without modifying model structure or routing mechanisms. The core innovation chains disk-resident expert weights through CPU-pinned memory staging buffers to GPU VRAM slots via asynchronous, non-blocking PCIe transfers. The system incorporates direct disk I/O via safetensors slicing (`MmapExpertStore`), intelligent pinned memory buffer management (`_next_staging()`) and an LRU cache with intra-batch freezing (`ExpertCache`) to prevent inadvertent eviction of required experts during batch processing. Empirical evaluation demonstrates that CacheFlow achieves **5–9 $\times$ reductions in peak GPU memory** compared to baseline systems across constrained capacity settings. Specifically, at capacity constraint $C = 1$ (single expert slot per MoE layer), the system maintains a peak VRAM footprint of **988.17 MB** while achieving per-token latency of **598.23 ms**, improving to **462.27 ms/token** at $C = 8$ (1845.26 MB VRAM)—a **23% latency reduction**—with throughput scaling from **100.62 to 130.04 tokens/minute**. The intra-batch freezing mechanism adds negligible overhead ($\sim$ 0.1% latency) while guaranteeing computational correctness. By enabling deployment of state-of-the-art MoE models on consumer-grade GPUs (RTX 3060, Jetson Orin with 12 GB VRAM) previously requiring specialized infrastructure, CacheFlow democratizes access to sparse expert architectures and establishes a practical foundation for scalable, efficient MoE inference in resource-constrained environments.

Keywords: Asynchronous expert prefetching, Edge device deployment, LRU cache with intra-batch freezing, Memory-efficient GPU scheduling, Mixture-of-Experts (MoE) inference, Resource-constrained inference optimization, Three-tier memory hierarchy.

Contents

Declaration	i
Supervisor’s Declaration	ii
Dedication	iii
Acknowledgment	iv
Ethical Statement	v
Table of Contents	v
List of Figures	vii
List of Tables	viii
1 Introduction	1
1.1 Background and Motivation	1
1.2 Problem Statement	2
1.3 Proposed Solution: CacheFlow	3
1.4 Research Objectives	3
1.5 Key Contributions	4
1.6 Thesis Organization	5
2 Literature Review	7
2.1 Overview	7
2.2 Foundational and Modern Architectures for Scalable Neural Models .	8
2.2.1 Adaptive Mixtures of Local Experts by Jacobs et al. (1991) .	8
2.2.2 Hierarchical Mixtures of Experts by Jordan and Jacobs (1994)	9
2.2.3 Adam: A Method for Stochastic Optimization by Kingma and	
Ba (2014)	9
2.2.4 Distilling the Knowledge in a Neural Network by Hinton et al.	
(2015)	9

2.2.5	Neural Machine Translation by Jointly Learning to Align and Translate by Bahdanau et al. (2015)	10
2.2.6	Batch Normalization by Ioffe and Szegedy (2015)	10
2.2.7	Layer Normalization by Ba et al. (2016)	10
2.2.8	Gaussian Error Linear Units (GELU) by Hendrycks and Gimpel (2016)	11
2.2.9	Attention Is All You Need by Vaswani et al. (2017)	11
2.2.10	Outrageously Large Neural Networks: The Sparsely-Gated Mixture-of-Experts Layer by Shazeer et al. (2017)	11
2.2.11	Efficient Softmax Approximation for GPUs by Grave et al. (2017)	12
2.2.12	Improving Language Understanding by Generative Pre-Training by Radford et al. (2018)	12
2.2.13	Breaking the Softmax Bottleneck by Yang et al. (2018)	12
2.2.14	Language Models are Unsupervised Multitask Learners by Radford et al. (2019)	13
2.2.15	Generating Long Sequences with Sparse Transformers by Child et al. (2019)	13
2.2.16	Root Mean Square Layer Normalization by Zhang and Sennrich (2019)	13
2.2.17	Language Models are Few-Shot Learners by Brown et al. (2020)	14
2.2.18	Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer by Raffel et al. (2020)	14
2.2.19	DeepSpeed: System Optimizations Enable Training Deep Learning Models with Over 100 Billion Parameters by Rasley et al. (2020)	14
2.2.20	An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale by Dosovitskiy et al. (2020)	15
2.2.21	GLU Variants Improve Transformer by Shazeer (2020)	15
2.2.22	Longformer: The Long-Document Transformer by Beltagy et al. (2020)	15
2.2.23	ALBERT: A Lite BERT for Self-Supervised Learning of Language Representations by Lan et al. (2020)	16
2.2.24	GShard: Scaling Giant Models with Conditional Computation and Automatic Sharding by Lepikhin et al. (2021)	16
2.2.25	Switch Transformers: Scaling to Trillion Parameter Models with Simple and Efficient Sparsity by Fedus et al. (2022)	16
2.2.26	FlashAttention: Fast and Memory-Efficient Exact Attention with IO-Awareness by Dao et al. (2022)	17

2.2.27	FlashAttention-2: Faster Attention with Better Parallelism and Work Partitioning by Dao (2023)	17
2.2.28	Efficient Memory Management for Large Language Model Serving with PagedAttention by Kwon et al. (2023)	17
2.2.29	LLaMA: Open and Efficient Foundation Language Models by Touvron et al. (2023)	18
2.2.30	SwapMoE: Serving Off-the-Shelf MoE-based Large Language Models with Tunable Memory Budget by Yoo and Lee (2023)	18
2.2.31	MoE-Infinity: Efficient MoE Inference on Personal Machines with Sparsity-Aware Expert Cache by Zuo et al. (2023)	19
2.2.32	DeepSpeed-MII: Model Inference Interface by Microsoft (2023)	19
2.2.33	TensorRT-LLM: Optimized Inference for Large Language Models by NVIDIA (2023)	19
2.2.34	DeepSeek-V2: A Strong, Economical, and Efficient Mixture-of-Experts Language Model by Liang et al. (2024)	20
2.2.35	fMoE: Fine-Grained and Prompt-Aware Expert Offloading for Large Mixture-of-Experts Serving by Shen et al. (2024)	20
2.3	Comparative Summary of Reviewed Systems	20
2.4	Research Gap	23
2.5	Summary	23
3	Methodology	25
3.1	Overview of the Proposed Architecture	25
3.2	System Architecture and Memory Pipeline	26
3.2.1	Three-Tier Memory Hierarchy	26
3.2.2	Direct Disk I/O via MmapExpertStore	27
3.2.3	Pinned Memory Staging and Async PCIe Transfer	28
3.2.4	ExpertCache: GPU-Resident Cache with Freezing	29
3.3	Mathematical Formulation of CacheFlow Routing	30
3.3.1	Memory Budget Decomposition with Three-Tier Architecture	30
3.3.2	Host Pinned Memory Footprint for Staging	31
3.3.3	Disk Storage and the Complete Memory Budget	31
3.3.4	Memory Reduction Formula	32
3.3.5	Dynamic Cache Buffer Definition	32
3.3.6	Token Routing and Cache Hit Efficiency	32
3.4	Asynchronous Execution Flow	33
3.4.1	Forward Pass Initialization and Batch Setup	33
3.4.2	Asynchronous Prefetching and Pending Transfers	34
3.4.3	Synchronization Before Expert Computation	35

3.4.4	Execution Ordering for Latency Minimization	35
3.4.5	Output Concatenation and Cache Unfreezing	36
3.5	Advanced LRU Cache Management and Freezing	36
3.5.1	Cache State Representation with Freezing	36
3.5.2	Cache Hit Path	37
3.5.3	Cache Miss with Available Free Slots	37
3.5.4	Cache Miss with Full Capacity (Eviction)	37
3.5.5	Complete Cache Lifecycle During Forward Pass	38
3.6	Implementation Environment	40
3.6.1	Software Stack	40
3.6.2	Hardware Configuration	40
3.6.3	Model Under Test	41
3.6.4	Experimental Protocol	41
3.6.5	Correctness Validation	43
3.7	Summary	43
4	Results and Discussion	45
4.1	Overview of Experimental Evaluation	45
4.2	Comparison with Existing Frameworks	45
4.3	Memory Efficiency vs. Computational Cost Trade-off	47
4.4	Cache Hit Rate and Asynchronous Routing Dynamics	50
4.5	Generalization Across Diverse MoE Architectures	52
4.6	Discussion and Limitations	53
4.6.1	Key Findings and System Insights	53
4.6.2	Limitations and Constraints	54
4.6.3	Future Directions and Improvements	54
4.7	Summary of Results	55
5	Conclusion	56
5.1	Summary of Research	56
5.2	Summary of Key Findings	57
5.2.1	Memory Efficiency	57
5.2.2	Latency and Throughput Performance	57
5.2.3	Cache Hit Rate and Temporal Locality Exploitation	58
5.2.4	Generalization Across MoE Architectures	59
5.2.5	Comparative Advantage Over State-of-the-Art Baselines	59
5.3	Limitations of the Study	59
5.3.1	I/O Bandwidth as a Fundamental Bottleneck	60
5.3.2	Single-GPU Inference Scope	60

5.3.3	Cache Hit Rate Saturation	61
5.3.4	Staging Buffer Minimalism and Saturation Effects	61
5.3.5	Scope Limitations on Evaluation	61
5.4	Future Work	62
5.4.1	Predictive Prefetching Based on Token Embeddings	62
5.4.2	Dynamic and Adaptive Staging Buffer Sizing	62
5.4.3	Multi-GPU Cache Coherency and Expert Parallelism	63
5.4.4	Token Reordering and Routing-Aware Batching	63
5.4.5	Hardware-Specific Optimizations	64
5.4.6	Theoretical Analysis of Cache Optimality	64
5.5	Concluding Remarks	64
	References	66

List of Figures

3.1	The three-tier CacheFlow memory hierarchy illustrates the complete data pathway from persistent disk storage through pinned CPU memory staging to GPU VRAM. Expert weights are stored compactly in safetensors format on disk. The MmapExpertStore reads specific expert slices on-demand. Pinned memory buffers stage the weights for asynchronous transfer. The GPU cache maintains the working set of experts. The diagram shows multiple experts at different stages: some resident in GPU cache, some pending transfer to GPU and some available for disk loading.	30
3.2	The asynchronous execution flow demonstrates the overlap of disk I/O, PCIe transfer and GPU computation. At Time T0, Expert A is being read from disk, Expert B is being transferred via PCIe and Expert C is being computed on the GPU. The pending dictionary tracks in-flight transfers. Cache freezing ensures required experts are protected. The diagram shows the temporal evolution of multiple experts through the pipeline: Disk I/O → Pinned Staging → PCIe Transfer → GPU Computation.	36
3.3	The advanced LRU cache with freezing mechanism shows how intra-batch protection works. Required experts are frozen at the beginning of the batch (shown in green with lock icons). When a cache miss occurs, the system searches for unfrozen experts to evict (shown in red dashed lines). Frozen experts cannot be evicted, ensuring their protection. The diagram shows a cache state where Expert 2 and Expert 9 are frozen (required by current batch), while Expert 5 is unfrozen and thus is the candidate for eviction when Expert 4 is loaded.	39
4.1	Peak VRAM Utilization Across Capacity Constraints. Peak GPU memory (in MB) scales monotonically with capacity constraint C . Multiple staging buffer configurations ($S = 1, 2, 4$) show overlapping curves, indicating that staging buffer count has negligible impact on GPU VRAM once the cache is the dominant memory consumer. . . .	49

4.2	Per-Token Inference Latency Across Capacity Constraints. Latency decreases from 598 ms/token at $C = 1$ to 462 ms/token at $C = 8$, representing a 23% improvement. The improvement rate slows beyond $C = 4$, indicating saturation.	49
4.3	Throughput (Tokens per Minute) Across Configurations. Throughput ranges from 96.8 to 130 tokens/min across capacity constraints. CacheFlow at $C = 8$ achieves 130.04 tok/min, approaching baseline systems.	50
4.4	Input Cache Hit Rate (%) as a Function of Capacity and Staging Buffers. Cache hit rates increase dramatically as capacity constraint increases from $C = 1$ to $C = 8$, rising from 9.03% to 35.32%. The heatmap reveals that staging buffer count has minimal impact on hit rates.	52
4.5	Inference Latency on Granite-3B (48 MB Experts) vs. Active Cache Capacity. Latency decreases from 483 ms/token at $C = 1$ to 298 ms/token at $C = 8$, a 38% improvement. The saturation beyond $C = 4$ indicates that cache capacity ceases to be the bottleneck.	53

List of Tables

2.1	Comparative Summary of Reviewed Architectures and Systems	22
4.1	Quantitative Comparison of CacheFlow with Existing Frameworks. Peak VRAM (MB), Throughput (tok/s), Latency (ms/token) and Hit Rate (%) across three configurations on the same model and hardware. CacheFlow demonstrates 3.5–9× reductions in GPU memory footprint while maintaining competitive or superior latency.	46
4.2	Peak VRAM, Latency, Throughput, and Hit Rate Across Capacity Constraints.	47
4.3	Input Cache Hit Rates (%) Across Capacity and Staging Configurations.	50
4.4	CacheFlow Performance on Diverse MoE Architectures.	52

Chapter 1

Introduction

1.1 Background and Motivation

The emergence of Mixture-of-Experts (MoE) architectures has marked a transformative shift in large-scale deep learning, enabling the training and deployment of increasingly capable language models within fixed computational budgets. Prominent examples include Mistral 7B (sparse expert switching with 32 experts), Qwen2-MoE-A2.7B and Qwen2-MoE-57B (with 64 experts per layer) and Mixtral 8x7B (eight expert pathways per token). These models achieve state-of-the-art performance on numerous benchmarks by routing each token to a small subset of specialized experts, thereby increasing model capacity while maintaining efficient computation.

The theoretical appeal of MoE models is well-established: they provide favorable compute-to-parameter ratios compared to dense alternatives. A 60-expert MoE layer processes tokens through only 2-4 selected experts per token, incurring $\mathcal{O}(\text{num_selected_experts})$ compute while maintaining $\mathcal{O}(\text{num_total_experts})$ parameters. However, this parameter abundance creates an acute practical challenge: the total weight matrices of all experts must be accessible during inference, consuming enormous quantities of GPU VRAM. A typical MoE layer with 60 experts, each weighing ~ 100 MB, requires 6 GB of expert parameters alone. When combined with transformer static parameters (embeddings, attention layers, output projections), KV caches for autoregressive decoding and intermediate activation buffers, the cumulative GPU memory demand can easily exceed 20-30 GB—far beyond the capacity of consumer-grade GPUs (typically 6-12 GB) and even many mid-range server accelerators.

This VRAM bottleneck has emerged as a critical barrier to the widespread deployment of MoE models. While the original MoE training and inference research often assumed access to highly specialized hardware with abundant memory, practical deployment scenarios frequently lack such resources. Academic researchers, small organizations and edge computing scenarios are forced to either abandon MoE models

entirely or resort to crude approximations (expert pruning, quantization) that sacrifice model quality. This motivates the development of efficient memory-aware inference systems that decouple GPU memory requirements from total expert parameter count, enabling deployment of full-capacity MoE models on resource-constrained hardware.

1.2 Problem Statement

Existing approaches to address the MoE memory bottleneck fall into two broad categories, each with significant limitations.

Distributed Expert Parallelism: Systems such as DeepSpeed-MoE and Megatron employ expert parallelism, where experts are partitioned across multiple GPUs with synchronized communication. While effective in datacenter environments with high-bandwidth interconnects, this approach introduces substantial complexity and is inaccessible to practitioners with single-GPU or limited hardware budgets. Furthermore, synchronous communication at expert boundaries creates hard latency floors that preclude deployment in latency-sensitive applications.

Host Memory Offloading (Naive Staging): An alternative strategy loads expert weights to host (CPU) pinned RAM prior to inference, then transfers them to GPU on-demand. This approach (exemplified by MoE-Infinity and other disk-offloading frameworks) reduces the GPU VRAM requirement but introduces a new bottleneck: the sequential nature of the prefill phase. When the system must load an expert from pinned host memory to GPU memory before computing on it, the forward pass stalls, blocking token processing until the PCIe transfer completes. Empirically, this results in latencies of **920 ms/token** or higher—a **1.5-3 $\times$ penalty** compared to baseline systems. Moreover, maintaining all experts in pinned host RAM incurs **5-10 GB** of host memory allocation, which constrains usable GPU capacity and limits the number of concurrent inference requests.

Joint Limitations: Both approaches fail to exploit a critical opportunity: **asynchronous overlapping of disk I/O, PCIe transfers and GPU computation**. In a synchronous design, the critical path is: (disk read) $\rightarrow$ (wait for completion) $\rightarrow$ (PCIe transfer) $\rightarrow$ (wait for completion) $\rightarrow$ (GPU compute). Each stage serializes, resulting in severe latency penalties. Furthermore, neither system maintains a principled caching strategy to exploit temporal locality in expert access patterns; experts are repeatedly transferred even if accessed consecutively.

The root cause of these limitations is architectural: existing systems do not implement a true three-tier memory hierarchy with non-blocking asynchronous operations. They either consolidate experts in GPU memory (scalability failure), consolidate experts in pinned host memory (latency failure), or implement synchronous staging (latency failure). No existing system seamlessly chains disk I/O $\rightarrow$ pinned staging $\rightarrow$

asynchronous GPU transfer with an intelligent caching layer.

1.3 Proposed Solution: CacheFlow

CacheFlow addresses these limitations through a comprehensive three-tier memory hierarchy and asynchronous execution model. The system is architected as follows:

Tier 1: Disk Storage (Persistent Archive): Expert weights are stored on disk in safetensors format, a compact binary encoding that enables efficient random-access slicing. Unlike naive approaches that read entire weight files, CacheFlow reads only the requested expert’s slice—typically 100-256 MB per expert, drastically reducing I/O volume.

Tier 2: Pinned CPU Memory (Staging Buffers): A small fixed-size pool of pinned (page-locked) CPU buffers stage weights during loading from disk. The key innovation is that the number of staging buffers (e.g., 4-8) is **decoupled** from the total number of experts (e.g., 60). This allows the system to pipeline multiple disk reads and PCIe transfers concurrently, exploiting the high I/O throughput available from modern storage devices and PCIe interconnects.

Tier 3: GPU VRAM (Working Set Cache): A fixed-capacity LRU cache in GPU memory holds the subset of experts currently in use. Critically, this cache employs an **intra-batch freezing mechanism**: at the beginning of each forward pass, all required experts (determined by token routing) are locked in the cache, guaranteeing they will not be evicted while the batch is computing. This eliminates correctness hazards and enables safe asynchronous loading of cold experts.

Asynchronous Execution Model: The system schedules disk I/O, PCIe transfers and GPU computation concurrently via `torch.cuda.Stream`. While Expert A’s weights are being read from disk by the CPU, Expert B’s previously-staged weights are simultaneously being transferred to GPU VRAM via PCIe and the GPU is computing on Expert C’s weights already resident in cache. This triple concurrency dramatically reduces effective latency.

Transparent to Application: CacheFlow does not modify the underlying MoE architecture, model structure, or routing mechanisms. It operates at the expert module level as a drop-in replacement for standard parallel expert implementations, making it broadly compatible with existing frameworks and models.

1.4 Research Objectives

This thesis aimed to achieve the following objectives:

- **Objective 1:** Design and implement a memory-efficient three-tier architecture

for MoE inference that decouples GPU VRAM requirements from total expert parameter count while maintaining computational correctness.

- **Objective 2:** Demonstrate that asynchronous prefetching via non-blocking PCIe streams can effectively overlap disk I/O, data staging and GPU computation, mitigating latency penalties typically incurred by disk offloading.
- **Objective 3:** Empirically validate the system across diverse capacity constraints (GPU memory budgets from 900 MB to 2 GB) and model architectures (Granite-3B, Qwen2-MoE variants), quantifying the achievable memory-latency trade-off frontier.
- **Objective 4:** Provide architectural insights and design guidelines for practitioners seeking to deploy large MoE models on resource-constrained hardware, including memory-optimal capacity configurations and staging buffer tuning strategies.

1.5 Key Contributions

The CacheFlow system and accompanying evaluation contribute the following advances to the field:

1. A Novel Three-Tier Memory Hierarchy with Asynchronous Staging:

CacheFlow introduces a principled architectural design that cleanly separates concerns across three memory tiers—disk, pinned host RAM and GPU VRAM—with asynchronous non-blocking data transfers. This design is fundamentally different from prior work and achieves several key properties: (a) disk I/O is proportional only to a single expert’s size, not the entire weight file; (b) pinned memory allocation is bounded by the number of staging buffers (typically 4-8), not the total expert count; (c) GPU cache capacity is independently tunable via the capacity constraint C without affecting host-level allocations.

2. Intra-Batch Freezing for Correctness and Safety:

The `freeze_many()` mechanism provides a simple yet powerful correctness guarantee: experts required during a batch cannot be evicted while the batch is computing. This eliminates deadlocks, data races and partial-result loss that would otherwise occur in naive asynchronous systems. The freezing protocol adds negligible overhead ($\sim 0.1\%$ latency) while providing transparent safety semantics.

3. Quantified Memory-Latency Trade-off Frontier:

Empirical evaluation demonstrates that CacheFlow achieves **5-9 $\times$ reductions in peak GPU VRAM** compared to baseline systems across constrained capacity settings. Specifically:

- At $C = 1$ (ultra-constrained, ~ 1 expert slot per MoE layer): **988 MB** peak VRAM, **598 ms/token** latency.
- At $C = 8$ (moderate constraint, ~ 8 expert slots per MoE layer): **1845 MB** peak VRAM, **462 ms/token** latency.
- Throughput maintained at **100-130 tokens/minute** across all configurations, compared to near-zero throughput for baseline systems in constrained settings.

This represents a practical Pareto frontier where practitioners can choose operating points balancing memory savings against acceptable latency overhead.

4. Architectural Generalization Across Diverse MoE Models:

CacheFlow’s design is model-agnostic and scales correctly across diverse MoE architectures. Evaluation on Granite-3B (8 experts, 48 MB each), primary test configuration (60 experts, 107 MB each) and analysis of Qwen2-MoE (64 experts, 256 MB each) demonstrates consistent performance characteristics. The system’s effectiveness is shown to improve with larger expert sizes, making it particularly valuable for next-generation large-scale MoE models.

5. Practical Deployment Enablement:

By reducing GPU memory requirements by an order of magnitude while maintaining reasonable latencies, CacheFlow enables deployment of production MoE models on consumer and mid-range GPUs. This democratizes access to state-of-the-art model architectures and reduces per-inference costs in datacenter environments through improved resource utilization.

1.6 Thesis Organization

The remainder of this thesis is structured as follows. **Chapter 2** provides a comprehensive literature review, discussing prior work in MoE model architectures, memory-efficient inference techniques and disk-offloading systems, contextualizing CacheFlow’s contributions within the broader research landscape. **Chapter 3** presents the complete methodological design of the CacheFlow system, detailing the three-tier memory hierarchy, the asynchronous execution model, the freezing mechanism and the mathematical formulation of memory budgets and capacity constraints. **Chapter 4** reports comprehensive experimental results, comparing CacheFlow against baseline systems (DeepSpeed MoE and MoE-Infinity), analyzing the memory-latency trade-off across varying capacity constraints, evaluating cache hit rates and prefetching dynamics and assessing generalization across diverse model architectures. **Chapter 5** concludes the thesis with a synthesis of key findings, discussion of system limitations and their

implications and identification of promising directions for future research. Throughout, emphasis is placed on rigorous quantitative evaluation, clarity of exposition and practical applicability to deployment scenarios.

Chapter 2

Literature Review

2.1 Overview

The modern large-scale models of language have developed based on the breakthroughs in terms of architecture, optimization, and systems design within the last three decades or so. Initial literature on Mixture-of-Experts (MoE) models proposed the use of conditional computation: a subsection of the model can specialize on specific areas of inputs. This was done with the goal of increasing the representational power of the network without corresponding increases in the cost of computation. Simultaneously, the development of optimization algorithms, normalization algorithms, and attention systems made the deep neural networks effective in scaling. These advances led to the development of Transformer-based architectures which were soon the standard design of sequence modeling. The more complex the models became in terms of millions of parameters, then billions of parameters, then even of trillions of parameters, the greater the increase in performance became in direct proportion to the corresponding rise in both computation and memory requirements.

Transformer changed the face of natural language processing. It rendered training very parallelisable and offered effective long -range dependency modeling. This was later followed up by more research that enhanced expressiveness, stability and scalability. Further innovations included as better normalisation techniques, new activation functions, sparse attention systems, and large scale distributed training systems. Although these advances increased the quality of models and the speed of training, they increased the level of inference to resource requirements. In dense ones, all of the parameters are required to be held in memory at inference, posing a difficulty in deploying the technology to a low-hardware device.

The response to the scaling problem was based on the sparse Mixture-of-Experts. MoE models incurred a massive number of parameters without proportionate raise in per-token compute by simply enabling a limited number of experts per-token.

However, as we have seen, inference systems do require all expert weights to be loaded on the GPU. As a result, the number of memory used tends to increase with the overall number of the experts, rather than the number of the experts at work. This restricts the capability of consumer GPUs to execute large MoE models.

Recent studies have changed the focus on inference time optimization and improved memory management. Efficient attention, key-value cache memory paging, task switching between the CPU and the GPU, and distributed sharding are some of the techniques that are expected to minimize memory bottlenecks. Although useful, these approaches are usually operated at a coarse granularity, not scheduling context of a specific system, and are oriented towards multi-gpu implementation instead of single-gpu environment. Therefore, there is still a research void of fine-grained, context-dependent expert scheduling on one single-processor graphics card.

In this chapter, a review of essential MoE literature, Transformer design, scaling techniques, and current inference systems takes the structure of a paper-by-paper. Both works are analyzed in terms of their idea, technical contribution, implementation, and empirical findings, as well as limitations, primarily in terms of inference time memory use. The history of the review: Mixture-of-Experts designs, the emergence of Transformers, large-scale sparse models, and current serving and memory-management designs. The chapter also concludes with a comparative summary and a definite research gap that drives CacheFlow.

2.2 Foundational and Modern Architectures for Scalable Neural Models

This part presents systematic, paper-by-paper discussions of traditional Mixture-of-Experts, approach to optimization, Transformer structures, large scale language model, and current inferential computations. The content of each subsection is a balanced paragraph, where one of the works presents its important idea, method, results, and limitations, especially concerning the inference-time memory scalability.

2.2.1 Adaptive Mixtures of Local Experts by Jacobs et al. (1991)

Jacobs et al. (1991) portrayed Adaptive Mixtures of Local Experts. Instead, they constructed a neural network that had several expert networks, each specializing in a segment of the input space. The process of a gating network, which is dynamically trying to assign weight to the output of each expert, is based on the soft partitioning process of probabilistic nature. An approximation of functionality in experiments was

higher than approximate converging monolithic networks. Nonetheless, the work concentrated on small-scale systems and it did not take into account memory limitations of hardware. The entire expert parameters were assumed to be loaded on inference, which upon the scale was a reasonable assumption but became an issue.

2.2.2 Hierarchical Mixtures of Experts by Jordan and Jacobs (1994)

Jordan and Jacobs (1994) built on the Mixture-of-Experts framework to introduce hierarchical gating structures which were optimized using Expectation-Maximisation. Davids would now have the ability to specialise at various levels, which would result in a higher degree of flexibility and capacity. Compared to flat mixtures the hierarchical design was better at classification and regression. As in the previous work, it also assumed all the experts were in memory at runtime and it did not concern runtime memory efficiency. However, the hierarchical concept also forms the basis of the contemporary sparse Transformer-based designs, according to which experts are conditionally activated, but usually reside in memory.

2.2.3 Adam: A Method for Stochastic Optimization by Kingma and Ba (2014)

Kingma and Ba (2014) presented Adam, which is an optimizer that combines adaptive learning rates and momentum to accelerate training. Adam follows first and second-order moments of gradients, stabilizing updates in large-scale training. It was made a default option of deep Transformers and MoE models due to its strength and little tuning. Adam on the other hand is only able to touch on the training but has no influence on the inference-time usage of memory. It assisted scaling training, but failed to alleviate the parameter-resident bottleneck that has arisen in implementing those trained models.

2.2.4 Distilling the Knowledge in a Neural Network by Hinton et al. (2015)

The authors Hinton et al. (2015) presented a method of knowledge distillation which is a more compact student model that has the accuracy of a larger neural network. The process takes the softened output probabilities of the teacher and transfers it to the student causing a number of parameters to be reduced in order to attain competitive performance. It has been experimentally verified that distilled models can be as accurate as larger networks by re-training, but rather than managing memory

dynamically during inference with a large model (which requires substantial memory), the compressed model gives a fixed map due to their minimal size. Therefore, distillation can only help in reducing memory use at the model level; it is less flexible with regards to expert-level memory management that is needed by sparse architectures.

2.2.5 Neural Machine Translation by Jointly Learning to Align and Translate by Bahdanau et al. (2015)

To simplify neural machine translation, Bahdanau et al. (2015) introduced attention to the decoding process to enable a model to pay attention to relevant input tokens. Their recurrent networks calculate weights on alignment at each decoding step and this enhances the quality of translation and long sequences are better handled. Attention however, raises the usage of intermediate memory due to the usage of alignment matrices and stored hidden states. Although this inspired Transformer, it does not cause the size of the parameter memory to be smaller; all the model weights are kept loaded to inference.

2.2.6 Batch Normalization by Ioffe and Szegedy (2015)

Ioffe and Szegedy (2015) suggested Batch normalization which stabilizes and accelerates the training process through normalization of layer inputs across mini-batches. It causes a lower covariate shift inside the model and permits the greater learning rates resulting in quicker convergence and higher generalization. However, Batch Normalization is essentially training friendly; it does not reduce the number of parameters and memory imprint of the model. In perception, the acquired scaling and shifting parameters remain functional, and the process does not foster the scalability of memory. Its primary contribution is performance with regard to trainings as opposed to the resource management on runtime mode.

2.2.7 Layer Normalization by Ba et al. (2016)

Two-dimensional normalization was introduced by Ba et al. (2016) as Layer Normalization that normalizes features activations, rather than batch ones. The method works better when batches statistics are volatile in recurrent and sequence models. Transformer architectures now incorporate Layer Normalization as a part of them. It increases training stability and convergence, and does not lower the demands of inference run time memory; during execution all the model parameters are resident in the GPU memory. It, therefore, adds architectural stabilization, rather than memory optimization.

2.2.8 Gaussian Error Linear Units (GELU) by Hendrycks and Gimpel (2016)

Hendrycks and Gimpel (2016) proposed a combination of ReLU and stochastic regularization, and called it the GELU activation function. GELU enhances the performance on the language and vision in Transformer feed-forward layers. Even while it has a positive impact on representational power, it does not alter the model size or memory footprint. Activation functions affect Churchill but not the need to load all weight matrices when making inferences. As such, its influence is on performance, rather than on the use of memory.

2.2.9 Attention Is All You Need by Vaswani et al. (2017)

The Transformer suggested by Vaswani et al. (2017) substitutes recurrent and convolutional sequence models with multi-headed self-attention and position-wise feed-forward networks. The architecture supports paralleling input token processing which enables high level of training efficiencies and scalability. Translate Experiments of machine translation demonstrate state-of-the-art by performance and lower training time than recurrent networks. Transformers however put all the parameters in memory at inference and therefore with increase in the model size, the memory requirements of the GPU also increase significantly. Although repetition bottlenecks are removed in the architecture, inference time memory scaling, which became important in large scale deployments was not addressed.

2.2.10 Outrageously Large Neural Networks: The Sparsely-Gated Mixture-of-Experts Layer by Shazeer et al. (2017)

Shazeer et al. (2017) proposed sparsely-gated Mixture-of-Experts layers, with the capacity to add parameter capacity but no expansion of per-token computation. Gating network chooses a small selection of experts per input whereas the experts allow conditional computation. It is shown that there is substantial improvement in language modeling and translation with computation manageable by experiments. However, not many experts are activated on a token, although the majority of systems maintain all expert parameters in memory, as part of inference. Thus, the use of memory depends on the number of total experts and not active ones. This separation of computational sparsity and memory density is one of the central points of later studies, such as CacheFlow.

2.2.11 Efficient Softmax Approximation for GPUs by Grave et al. (2017)

Grave et al. (2017) suggested effective softmax approximations in order to accelerate large-vocabulary language modeling on GPUs. They calculate the hierarchical clustering and adaptive partitioning method to minimize the computational complexity in the output layers. The training and inference time in experiments is faster than in standard softmax. Nevertheless, although the method is less costly in arithmetic, it does not reduce the requirement commonly required by all model parameters being resident at inference time. The focus of optimization is output-layer computation instead of the memory footprint and therefore is not concerned with the issue of weight residency in large Transformer or MoE models.

2.2.12 Improving Language Understanding by Generative Pre-Training by Radford et al. (2018)

The article by Radford et al. (2018) presented a new method of NLP in generative pre-training followed by fine-tuning as a potent tool. They use large language models based on transformers, which they trained on large unlabeled corpora, and distilled to downstream with minimal learning. The tests revealed significant gains in such benchmarking as question answering and textual entailment. The designs are based on dense Transformer layers, where all the parameters are opened during inference. The model supported the tendency of increasing the performance with the number of parameters, although it was smaller than subsequent large systems. It greatly developed transfer learning in NLP but failed to take into account memory limitations when running at runtime and how to manage the parameters dynamically.

2.2.13 Breaking the Softmax Bottleneck by Yang et al. (2018)

Yang et al. (2018) discovered the so-called softmax bottleneck, the constraint of expressiveness of typical softmax layers in language models. They put forward a high rank factorization that upscales representational capacity but at a minimal computational cost. On language-modeling benchmarks, it was found that there was a decreased perplexity. The approach, though, does not decrease the number of stored parameters or modify assumptions about memory residency: all of the learned matrices have to be available in order to perform inference. In turn, inference-time memory pressure, particularly that of large Transformer models, does not decrease, as modeling capacity is improved.

2.2.14 Language Models are Unsupervised Multitask Learners by Radford et al. (2019)

GPT-2 was presented by Radford et al. (2019), who showed that the zero-shot setting allows large-scale language models which have been trained on a variety of web data to handle a number of tasks. The model takes a deep Transformer decoder that has much more parameters than previous systems. The findings indicate high performance in translation, summary and question answering without fine-tuning. Yet it is a completely dense architecture, and must have full parameter residency on inference. The size of the model (billions of parameters) renders the use of GPUs memory intensive, which tends to increase deployment costs due to memory limits.

2.2.15 Generating Long Sequences with Sparse Transformers by Child et al. (2019)

To reduce the quadratic cost of self-attention, Child et al. (2019) presented Sparse Transformers. They encode orderly sparsity in attention matrix that allows the sequential processing of long sequences at a lower rate of memory and computation. Better results in experiments on long context language models are demonstrated. This low density reimbursement attention-related memory artifact, but not residency of model weights. In inference, all the feedforward and projection parameters are loaded in their entirety. Therefore, attention sparsity helps in alleviating sequence-length scaling, but does not tackle the bigger problem of expert weight memory in large MoE systems.

2.2.16 Root Mean Square Layer Normalization by Zhang and Sennrich (2019)

Zhang and Sennrich (2019) suggested a simplified version of Layer Normalization called Root Mean Square Layer Normalization (RMSNorm). As computational efficiency and numerical stability RMSNorm is a variant of mean-centered variance-normalized normalization. Tests have demonstrated performance similar and better than standard Layer Normalization in Transformers. RMSNorm is a computationally less costly but does not modify the overall number of the parameters and the memory residency of the parameters in the inference. It makes its contribution in optimization and stability rather than scalability of the memory usage. Even large Transformer and MoE models based on RMSNorm require full weight residency.

2.2.17 Language Models are Few-Shot Learners by Brown et al. (2020)

GPT architecture Transformers were expanded to 175 billion parameters, presented by Brown et al. (2020). The primary contribution is that the demonstration of powerful few-shot and zero-shot learning without either fine-tuning or scaled down to huge sizes. The dense decoder is trained using huge datasets using distributed infrastructure. Findings depict significant improvements on NLP benchmarks (question answering and translation). The dense structure however requires full parameter residency to be applied in inference, and therefore cannot be deployed locally with specialized clusters. Although GPT -3 confirms scaling laws with respect to performance, it illustrates drastic limitations of dense large-scale models in memory.

2.2.18 Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer by Raffel et al. (2020)

Raffel et al. (2020) suggested T5, which resolves all NLP tasks to a single text-to-text model. The structure is built based on a standard Transformer encoder-decoder that is trained on a huge corpus that is under curation. It shows high transfer-learning on various benchmarks. Although T5 enhances training methodology and task unification, the subnet has a dense execution of parameters that necessitate an entire model reservoir throughout the inference phase. When T5 is scaled up, the amount of memory required by the GPU is scaled up accordingly. The work makes progress on modeling generality but no progress on inference time memory constraints.

2.2.19 DeepSpeed: System Optimizations Enable Training Deep Learning Models with Over 100 Billion Parameters by Rasley et al. (2020)

Rasley et al. (2020), proposed system-level DeepSpeed, which is an efficient training of extremely large models by optimization of memory and distributed partitioning. The ZeRO (Zero Redundancy Optimizer) of DeepSpeed divides optimizer states and gradients in the devices which allows less memory duplication through training. Experiments demonstrate that it is possible to train models of more than 100 billion parameters on commodity GPU clusters. DeepSpeed is however mainly aimed at training efficiency and distributed scaling, and not optimization of single node inference. Dense models even during inference still need full parameter residency, unless they are specifically partitioned. DeepSpeed cuts, therefore, reduce but not eliminate, the redundancy of memory accessible to of-the-shelf inference on a single gpu.

2.2.20 An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale by Dosovitskiy et al. (2020)

Dosovitskiy et al., (2020) proposed the Vision Transformer (ViT) that uses the Transformer architecture to deal with image recognition by splitting an image into patches that can be regarded as tokens. The model is able to equal convolutional neural networks on massive datasets, interface attention-based designs outside of NLP. Similarly to language Transformers, ViTs use dense parameter network models which also require full memory residency during inference. The larger model depth and width the more memory memory of a graphics card also require. Although the work extends the use of Transformer, it does not change the memory scaling properties.

2.2.21 GLU Variants Improve Transformer by Shazeer (2020)

Gated Linear Unit (GLU) variants were presented by Shazeer (2020) as variants of Transformer feed-forward layers focused on increasing their expressiveness. These variants introduce gating mechanisms to linear transformations and hence the variants are able to better skilled modeling capacity and enhance training stability. Experiments indicate that GLU versions do better with a variety of language-modeling tasks than do the normal feed-forward layers. The changes also, on some parameters, increase computational load and parameter count. Notably, there is no net reduction in inference-time memory GLU does not require the use of memory to decrease inference-time Memories: all transformed feed-forward parameters have to remain in memory during model execution, which is a dense-memory assumption.

2.2.22 Longformer: The Long-Document Transformer by Beltagy et al. (2020)

Proposed by Beltagy et al. (2020), the Longformer is the Transformer dispensing the entire idea of self-attention with a combination of local windowed attention and sparse global attention. The method reduces the quadratic memory and calculate cost of attention on long documents. Longformer has also been experimentally found to scale to long-sequence tasks, such as document classification and question answering, much better. Longformer is more economical in terms of attention-buffer memory, but the fact remains that not only do all model weights remain loaded when making inferences, but they are also held there. Attributes that represent the attention activations still take a backseat compared to the parameters in terms of occupying the GPU memory footprint.

2.2.23 ALBERT: A Lite BERT for Self-Supervised Learning of Language Representations by Lan et al. (2020)

Lan et al. (2020) proposed ALBERT, a downsized version of BERT, which downsizes model size without degrading performance through sharing the weights across the layers and factorizing the embedding matrices. These transformations decrease significantly the number of parameters total as compared to BERT. It has been experimentally demonstrated that ALBERT obtains competitive accuracy with less memory consumption. Nevertheless, it is the architectural compression, rather than the dynamic run time control, that is remembered. ALBERT reduces the number of parameters which are always static but does not have adaptive memory scheduling or conditional expert activation. As such it is a compression method as opposed to a runtime inference optimizer.

2.2.24 GShard: Scaling Giant Models with Conditional Computation and Automatic Sharding by Lepikhin et al. (2021)

Lepikhin et al. (2021) introduced GShard, which is a system that scales massive models with conditional computation as well as automatic sharding of parameters introduced in distributed hardware. It uses sparse layers of Mixture-of-Experts on Transformers and separates experts among modified devices so as to balance memory and computing duties. Experiments indicate that GShard has the ability to train very large MoE models on distributed clusters of TPUs. But the method is directed towards multi-device situations that already possess distributed memory. It makes no effort inference-time memory constraints on one GPU. Special weights remain full in the partition distributed and soft runtime scheduling at hard memory constraints is not investigated.

2.2.25 Switch Transformers: Scaling to Trillion Parameter Models with Simple and Efficient Sparsity by Fedus et al. (2022)

Fedus et al. (2022) proposed the Switch Transformer, that is, Mixture-of-experts routing is simplified through only one expert per token (top-1 routing) rather than using many experts. This minimizes the routing overhead and enhances stability of the training, enabling models to scale to trillion parameters. Experimental optimization has shown that Switch Transformers are equal to dense models in terms of cost per token. Nevertheless, the inference systems usually maintain all expert parameters

in memory to support dynamic routing, and therefore the memory consumption is directly proportional to the number of experts as opposed to the quantity of experts that are actually active. In this way, even though Switch Transformers can promote sparse scaling, they have no implication of automatically addressing inference-time memory residency issues.

2.2.26 FlashAttention: Fast and Memory-Efficient Exact Attention with IO-Awareness by Dao et al. (2022)

FlashAttention is an algorithm demonstrated by Dao et al. (2022) as it reduces memory overhead and accelerates self-attention. Focusing on memory-access patterns and being IO-aware, FlashAttention makes very precise attention calculations and at a far smaller memory footprint. Standard experiments demonstrate an increased throughput and low-memory consumption associated with processing long sequences in contrast to the usual attention. However, FlashAttention is concerned with the activation memory rather than the model parameters that are static. If it does not lower the need to have all the model weights such as expert weights in memory. Thus, FlashAttention can ease scale-to-sequence but is not a more powerful way of overcoming the expert-weight bottleneck in memory.

2.2.27 FlashAttention-2: Faster Attention with Better Parallelism and Work Partitioning by Dao (2023)

FlashAttention was further improved to a more-parallel organization and workload partitioning by Dao (2023) to create FlashAttention 2. It actively serving device is increasing the usage of the graphics card, reducing synchronization costs, which translates to additional performance and memory-saving improvements over the base. Similarly to FlashAttention, FlashAttention-2 does not focus on the parameter storage of Transformers or MoE models but instead, intermediate activation memory and compute efficiency. It therefore enhances the performance of the process of attention execution but fails to address the memory residency at the expertise level among sparse architectures.

2.2.28 Efficient Memory Management for Large Language Model Serving with PagedAttention by Kwon et al. (2023)

Kwon et al. (2023) proposed the memory-based attention method, which is PagedAttention, which can improve the efficiency of large-language-model serving. The

technique virtualizes key-value (KV) cache storage using paging of memory instead of consecutive blocks that make it be less prone to fragmentation and better utilize the GPU. Experiments demonstrate that memory waste and throughput are less in the case of long-context inference. Yet, PagedAttention only deals with the memory in terms of states of attention, as opposed to the model parameters. It is not a solution to maintaining all model weights, and, in particular, expert weights in MoE models, in GPU memory. In this way, the technique is a more efficient way of serving, but does not eliminate expert-level memory bottlenecks.

2.2.29 LLaMA: Open and Efficient Foundation Language Models by Touvron et al. (2023)

Touvron et al. (2023) presented a family of open and relatively efficient dense Transformer-based models, the use of which uses optimized strategies of data curation to propose LLaMA. The researchers demonstrated that intelligently developed training pipelines could be used to perform equally well with respect to larger systems using smaller models. Results of the experiment were good on various NLP benchmarks. Nevertheless, LLaMA is a large model that loads all the parameters in memory when performing inference. It also scales its memory requirement with the size of models, limiting its use to consumer GPUs. Thereby, although LLaMA is more efficient in terms of parameter optimization compared to the previous dense models, it lacks the capability to handle dynamic memory management.

2.2.30 SwapMoE: Serving Off-the-Shelf MoE-based Large Language Models with Tunable Memory Budget by Yoo and Lee (2023)

SwapMoE, suggested by Yoo and Lee (2023), is a system tapping expert weights in both CPU and GPU to allow the inference of large MoE models and limited memory. It has a memory limit fixed by the user and automatically transfers experts at on-demand. It was found that experiments indicated lower peak GPU memory and inference with stricter VRAM limits. Swapping is based upon heuristic/LRU-based policy and neither exploits fine-grained context nor any key-value reuse data. The round-trip transfer of CPU to and from GPUs may introduce a latency and stability may be compromised under intermittent loads. SwapMoE deals somewhat with expert memory residency and does not provide very advanced scheduling to reduce thrashing.

2.2.31 MoE-Infinity: Efficient MoE Inference on Personal Machines with Sparsity-Aware Expert Cache by Zuo et al. (2023)

The article by Zuo et al. (2023) proposes MoE-Infinity that leverages expert caching based on the idea of sparsity to execute MoE inferences on personal computers. The system evicts and loads experts depending on patterns of access, reducing the requirements of the GPU memories. Tests demonstrated that MoE models in large scale could be executed on commodity equipment which had managed memory consumption. The defaulting cache is based on basic eviction brush strokes, and does not utilize any token level context or key-value reuse indications. There could be communication overhead in the event of frequent repeated transfers between memory tiers. MoE-Infinity builds upon memory-aware MoE inference, but does not give the scheduling logic the continuity as a forward path of the model.

2.2.32 DeepSpeed-MII: Model Inference Interface by Microsoft (2023)

DeepSpeed-MII offers an inference interface that runs on DeepSpeed distributed optimization infrastructure, that allows serving large-model applications simultaneously using multiple GPUs. The system is based on model parallelism and pipeline using optimized inference to enhance throughput. It is not aimed at fine-grained expert-limited node level scheduling of MoE architectures, but at deploying clusters on a large scale fulfilling single-node memory constraints. Dense parameter residency is the default. DeepSpeed-MII advances the inference scale in a multi-device design, but is not concerned with the expertise weight memory management of restricted GPUs.

2.2.33 TensorRT-LLM: Optimized Inference for Large Language Models by NVIDIA (2023)

TensorRT-LLM is a highly optimized inference engine bringing the use of kernel fusion, optimization of precision, and hardware acceleration to improve the performance of large language models. It minimizes the latency and maximizes throughput by the way of low-level optimization and efficient utilization of GPUs. Choose any model However, TensorRT-LLM supposes that all model weights are resident in AV RAM prior to execution. It does not have active expert scheduling or weight-switching of MoE architectures. Although it accelerates the computation process, memory footprint is still associated with size of total parameters. In this way, TensorRT-LLM

is maximizing speed, without addressing the main memory scale issue of large systems using MoE.

2.2.34 DeepSeek-V2: A Strong, Economical, and Efficient Mixture-of-Experts Language Model by Liang et al. (2024)

Liang et al. (2024) came up with DeepSeek-V2, a massive Mixture-of-Experts language model that is more efficient in the use of parameters but equally does not compromise its performance. The architecture amalgamates distant expert layers to the Transformer backbone and concentrates on economical scaling in order to reduce the computational cost. Competitive accuracy was demonstrated to be better than dense models of approximately similar effective capacity and greater compute efficiency per token. But inference typically stores all expert weights in which to run dynamic routing are stored in the GPU memory. Therefore, DeepSeek-V2, however, makes scaling efficient, but fails to address the problem of memory residency on limited hardware.

2.2.35 fMoE: Fine-Grained and Prompt-Aware Expert Offloading for Large Mixture-of-Experts Serving by Shen et al. (2024)

Shen et al. (2024) developed fine-grained, prompt-sensitive expert offloading fMoE, an expert offloading mechanism to help to serve the needs of large Mixture-of-Experts model. The system is used to predict expert execution and purges inactive experts to identify prompt features, minimizing the peak memory in the GPU. It was experimentally demonstrated that memory efficiency was enhanced and fewer swaps were occurring compared to the case with the use of a static offloading. fMoE targets distributed serving environments and balances performance across devices. Even though it implements prompt-awareness, it does not schedule on a single GPU such that it is a strict part of the forward path of the model. Coordination overhead and routing analysis can be imposed by high concurrency on latency; fMoE is an expert aware offloading technology that does not consider context offloading with strong single node memory constraints.

2.3 Comparative Summary of Reviewed Systems

Mixture-of-Experts, dense Transformers, scaling innovations, and inference-serving systems are included in the literature. This is despite the fact that these studies enhanced modeling, efficiency and scalability but their approaches to its management

during the inference process vary widely. Very limited amount of systems are concerned with the maintenance of expert weights on the GPU in the event of limited memory. The following table gives the sparsity, expert granularity, memory strategy, deployment focus and inference-time memory constraints of each system.

Analytical Interpretation

The most important patterns in the analysis:

- **Dense Architectures and Static Residency:**

Most Transformer-based models store all the parameters in the GPU memory during inference. The weights remain resident, although they can be made cheaper to train, or they can be made much smaller during activations, preventing scalability.

- **Sparse Compute vs. Sparse Memory:**

The models of MoE minimize the calculation by utilizing a small number of experts on each token. However, in most instances, all the experts remain loaded even when not in use hence, memory sparsity is nonexistent.

- **Offloading-Based Solutions:**

SwapMoE and MoE-Infinity reduce peak VRAM memory by swapping of both the GPU and CPU memory. Such schemes are based on heuristics or LRU eviction, and may bring in transfer delays.

- **Distributed-Centric Designs:**

GShard and DeepSpeed are concerned with the optimization of multi-device partitions, as opposed to single-GPU. They also assist in groups but not in a single constrained node.

- **Limited Context-Aware Scheduling:**

Very little systems take advantage of token-level context or key-value reuse to make decisions on which experts to retain. The Majority make use of static rules or generic caches.

On the whole, the community came a long way in terms of the computational power and scaling, yet there is still nothing about finer/grained/context-aware work-scheduling on a single-GPU deployment.

Table 2.1: Comparative Summary of Reviewed Architectures and Systems

Related Work	Sparse Compute	MoE Layer	Expert Scheduling	CPU-GPU Swap	Consumer GPU Ready	Memory Limitation
Adaptive MoE (1991)	Yes	Yes	No	No	N/A	No runtime memory management considered.
Transformer (2017)	No	No	No	No	Yes	Full-weight residency required at inference.
Sparsely-Gated MoE (2017)	Yes	Yes	No	No	No	Experts typically remain fully resident.
Switch Transformer (2022)	Yes	Yes	No	No	No	Sparse compute; memory dominated by expert weights.
FlashAttention (2022)	No	No	No	No	Yes	Optimizes attention activations, not weight residency.
PagedAttention (2023)	No	No	No	No	Yes	KV-cache paging only; expert weights unchanged.
SwapMoE (2023)	Yes	Yes	Partial	Yes	Partial	Transfer overhead and thrashing under frequent swaps.
MoE-Infinity (2023)	Yes	Yes	Partial	Yes	Partial	Heuristic caching; limited context prioritization.
DeepSpeed-MII (2023)	No	No	No	Partial	No	Distributed focus; not single-GPU expert scheduling.
TensorRT-LLM (2023)	No	No	No	No	No	Assumes weights already resident in GPU memory.
DeepSeek-V2 (2024)	Yes	Yes	No	No	No	Does not inherently solve constrained expert residency.
fMoE (2024)	Yes	Yes	Partial	Yes	Partial	Designed for multi-device serving; coordination overhead.

2.4 Research Gap

Architecture (above), efficiency, and scaling The above literature review demonstrates advancements in architecture, performance, and size. Conditional computation was introduced by Mixture-of-Experts in order to increase model capacity without proportionately increasing compute. Parallel sequence modeling would then be made possible by transformers, and very large models such as GPT-3 had proved these scaling laws. Recent research addresses sparse routing, sharding, memory-efficient attention, and expert offloading to reduce compute and memory bottlenecks.

Nevertheless, a number of gaps are left open to inference of the limited GPU memory:

- Dense models necessitate the weights taking up all the weight, which is not feasible on consumer-grade devices.
- MoE continues to maintain all the experts loaded and thus memory is a function of the number of total experts and not the active experts.
- The offloading systems are coarse-grained, heuristic driven systems; they might provoke the frequent transfers of data, leading to latency and thrashing, and are oriented to the distributed systems.
- Attention optimizations do not just decrease the memory of the activations, but they primarily do not decrease the storage of the expert weights.

The literature taken together indicate a discrepancy: memory sparsity is not accompanied with arithmetic sparsity. Only a small number of systems take into account fine grained and context sensitive scheduling during the forward pass on one GPU with hard constraints. Expert weights are so far treated by few frameworks and ranked in order by reuse patterns that they can maintain their execution without architecture modification. Such a gap is the inspiration of memory aware scheduling of MoE inference in CacheFlow.

2.5 Summary

Chapter 2 provided a summary of scalable neural models, starting with ancient Mixture-of-Experts to contemporary inference serving models. Mixture-of-Experts This paper presented conditional computation; Transformers presented parallel sequence modeling, scaling tools such as sparse routing and sharding had increased

compute efficiency, but could do nothing about creating inference-time memory constraints. Subsequent literature covered compression, parameter sharing, efficient softmax and activation optimizations. Newer systems introduced memory-aware algorithms, such as KV-cache paging, CPU-GPU swapping, partitioning, although most of them still assume complete expert residency.

Computational sparsity in sparse MoE is not associated with memory sparsity, limiting the memory of GPUs. Those systems are primarily based on dense model acceleration, scaling distribution, or heuristically conducted offloading. Very limited support fine-grained, context-driven scheduling into a single-gpu execution path. This continuous gap leads to the requirement of CacheFlow, which is a framework that dynamically supports expert presence, according to run-time requirement and ensures proper execution.

The second chapter is based on this premise by stating the proposed system architecture and temporal scheduling approach of CacheFlow that would resolve the mentioned limitations of inference-time memory management of Mixture-of-Experts models.

Chapter 3

Methodology

3.1 Overview of the Proposed Architecture

The CacheFlow system is designed to enable efficient Mixture-of-Experts (MoE) inference on GPU hardware with severe memory constraints. The core innovation extends beyond simple GPU caching: it implements a comprehensive three-tier memory hierarchy that seamlessly chains disk-resident expert weights through pinned CPU RAM staging to GPU VRAM slots via asynchronous, non-blocking PCIe transfers. This architecture fundamentally decouples GPU memory requirements from the total expert parameter count, enabling deployment of massive MoE models on resource-constrained hardware.

The proposed approach is a system-level optimization that does not modify the underlying MoE architecture or transformer model structure. Instead, it operates at the layer level through the replacement of standard parallel expert modules with an optimized implementation that incorporates direct disk I/O, intelligent memory staging, asynchronous prefetching and temporal cache freezing mechanisms. Rather than maintaining all expert weights in host RAM (as in traditional approaches), CacheFlow leverages the observation that expert weight files are compactly stored on disk in the safetensors format. The system reads only the required expert weights directly from disk on-demand, stages them through CPU-pinned memory buffers and transfers them asynchronously to GPU slots. A sophisticated caching mechanism with intra-batch freezing ensures that frequently-accessed experts remain resident in GPU memory while allowing cold experts to be loaded and accommodated within a fixed capacity constraint.

This three-tier architecture—Disk, Pinned RAM and GPU VRAM—creates a memory-performance tradeoff landscape where capacity can be tuned per-application. The system maintains computational correctness while achieving dramatic reductions in GPU memory footprint and enabling models previously inaccessible to resource-

constrained environments.

3.2 System Architecture and Memory Pipeline

3.2.1 Three-Tier Memory Hierarchy

The CacheFlow system explicitly organizes expert weight storage and access across three distinct memory tiers:

- **Tier 1: Disk Storage (Persistent):**
 - Authoritative storage medium for all expert weight matrices.
 - Expert weights are persisted in `safetensors` format, a compressed binary serialization enabling efficient random access.
 - Access is performed via `MmapExpertStore`, which implements zero-copy slicing of expert weight tensors.
 - Large capacity (model-dependent, typically 10–100 GB for large MoE models).
 - Slow access bandwidth (approximately 100–500 MB/s depending on storage medium).
 - No GPU involvement; CPU reads weights directly from disk.
- **Tier 2: Pinned Host Memory (Staging Buffer):**
 - Fixed-size set of CPU-resident memory buffers allocated with CUDA page-locking semantics.
 - Serves as a temporary staging area for weights being loaded from disk.
 - Each staging buffer holds a single expert’s weight matrix.
 - Pinned allocation allows asynchronous, zero-copy PCIe transfers to GPU without CPU intervention.
 - Typically allocated with `torch.cuda.pin_memory()` or platform-specific pinning APIs.
 - Number of staging buffers (`staging_slots`) is a configurable hyperparameter (e.g., 8 or 16).
 - Acts as a decoupling point between slow disk I/O and fast GPU transfer.
- **Tier 3: GPU VRAM (Working Set):**
 - Reserved for active computation and inference.

- Contains: model static parameters (embeddings, attention layers, output head), intermediate activations and a fixed set of expert weight slots in GPU memory.
- Limited capacity typically between 6 GB and 24 GB on consumer-grade hardware.
- Ultra-high-bandwidth access (up to 900 GB/s on modern NVIDIA GPUs).
- Expert slots are implemented as a preallocated `torch.Tensor` of shape `(num_gpu_slots, output_size, input_size)`.

3.2.2 Direct Disk I/O via MmapExpertStore

The `MmapExpertStore` class abstracts expert weight access directly from disk storage. Unlike traditional approaches that load entire weight matrices into RAM, `MmapExpertStore` implements memory-efficient slicing:

Initialization Phase:

```
store = MmapExpertStore(weight_file)
```

The constructor calls `_read_metadata()`, which inspects the safetensors file header to determine:

- The tensor key associated with expert weights (typically "model.experts.weight" or equivalent).
- The tensor shape `(num_experts, output_size, input_size)`.
- The tensor data type (float32, float16, or int8).

This metadata reading is $O(1)$ in complexity and requires only header inspection; no weight data is loaded into memory.

Per-Expert Loading:

```
expert_weight = store.load_expert(expert_id, out=staging_buffer)
```

The `load_expert()` method executes the following steps:

1. **Bounds Checking:** Verify that $0 \leq \text{expert_id} < \text{num_experts}$.
2. **Safetensors Slicing:** Open the safetensors file and invoke `get_slice()` to extract only the weight matrix corresponding to `expert_id`. Safetensors' slice operation reads precisely `output_size × input_size × dtype_bytes` bytes from disk, skipping the preceding experts' data.

3. **Buffer Population:** Copy the sliced weights into the provided out buffer (a pinned memory allocation). This copy operation uses efficient CPU memory operations but does not transfer data to GPU at this stage.
4. **Return:** Return the populated buffer tensor.

The key advantage is that disk I/O is proportional to a single expert's size, not the entire weight file. For a model with 60 experts of 100 MB each (6 GB total), loading one expert reads only 100 MB from disk.

3.2.3 Pinned Memory Staging and Async PCIe Transfer

The `_next_staging()` method in `GraniteMoeParallelExperts` manages the lifecycle of pinned memory buffers:

```
def _next_staging(self) -> Optional[torch.Tensor]:
    # Returns an available pinned buffer, or waits for
    # one to be free.
    # After the buffer is used for disk I/O, it is
    # immediately scheduled for asynchronous transfer to
    # GPU. The system waits for prior transfers if needed.
```

The staging pipeline operates as follows:

1. **Buffer Availability Check:** Maintain a list of pinned buffers. If all buffers are currently in-flight (pending PCIe transfers), the system blocks until one completes.
2. **Synchronization Point:** Before reusing a pinned buffer, ensure that any prior PCIe transfer using that buffer has completed. This is tracked via `torch.cuda.Event` objects.
3. **Return Buffer:** Return an available pinned buffer for the next disk read operation.

After a weight is loaded from disk into a pinned buffer, an asynchronous PCIe transfer is scheduled using `torch.cuda.Stream`:

```
stream = torch.cuda.Stream()
with torch.cuda.stream(stream):
    gpu_slot[:] = pinned_buffer # Copy via PCIe
    event = torch.cuda.Event()
    event.record(stream) # Record completion in stream
pending[expert_id] = (event, gpu_slot) # Track pending transfer
```

This design allows multiple experts' weights to be loaded from disk and transferred to GPU concurrently. While Expert A's weights are being read from disk by the CPU, Expert B's previously-staged weights are being transferred to GPU via PCIe and Expert C is being computed on the GPU. This overlapping of I/O, transfer and computation dramatically reduces overall latency.

3.2.4 ExpertCache: GPU-Resident Cache with Freezing

The `ExpertCache` class manages a fixed-size LRU cache of experts in GPU memory. Beyond standard LRU operations, it introduces a critical freezing mechanism to prevent intra-batch conflicts:

Cache State:

- `cache: OrderedDict[int, torch.Tensor]`: Maps expert IDs to their GPU-resident weight tensors.
- `_frozen: set[int]`: Set of expert IDs that are protected from eviction (locked during current batch processing).
- `capacity: int`: Maximum number of experts that can be cached simultaneously.

Freezing Mechanism:

```
def freeze_many(self, expert_ids: List[int]) -> None:
    self._frozen.update(expert_ids)

def unfreeze_many(self, expert_ids: List[int]) -> None:
    for expert_id in expert_ids:
        self._frozen.discard(expert_id)
```

At the beginning of a forward pass, the set of required experts (determined by token routing) is frozen in the cache. This ensures that if the system must fetch an expert that is not cached (a cold expert), it cannot accidentally evict a required expert that is already cached. Freezing is released at the end of the batch, allowing normal LRU eviction to proceed for the next batch.




Figure 3.1: The three-tier CacheFlow memory hierarchy illustrates the complete data pathway from persistent disk storage through pinned CPU memory staging to GPU VRAM. Expert weights are stored compactly in safetensors format on disk. The MmapExpertStore reads specific expert slices on-demand. Pinned memory buffers stage the weights for asynchronous transfer. The GPU cache maintains the working set of experts. The diagram shows multiple experts at different stages: some resident in GPU cache, some pending transfer to GPU and some available for disk loading.

3.3 Mathematical Formulation of CacheFlow Routing

3.3.1 Memory Budget Decomposition with Three-Tier Architecture

The total GPU memory consumption for an optimized MoE model is decomposed across the three-tier hierarchy:

$$M_{VRAM} = M_{static} + M_{cache} + M_{KV} + M_{act} \quad (3.1)$$

Where:

- M_{static} : Memory consumed by static model parameters (embeddings, attention layers, output head, layer normalization) that remain permanently in GPU memory.
- M_{cache} : Memory allocated to expert weight slots in GPU VRAM (the working set).
- M_{KV} : Memory for key-value caches in attention mechanisms (if applicable).
- M_{act} : Memory for intermediate activations during forward computation.

The critical innovation is that M_{cache} is now **independently constrained** from the total expert parameter count. The cache memory is governed by:

$$M_{cache} = L_{MoE} \times C \times P_{expert} \times b_o \quad (3.2)$$

Where:

- L_{MoE} : Number of MoE layers in the model.
- C : Capacity constraint—the maximum number of expert slots in GPU VRAM per layer (e.g., $C = 4$ means at most 4 experts reside in GPU VRAM per layer simultaneously).
- P_{expert} : The size of a single expert weight matrix (in number of parameters).
- b_o : The byte size per parameter (typically 2 bytes for float16, 4 bytes for float32).

3.3.2 Host Pinned Memory Footprint for Staging

The pinned memory staging buffer allocation is distinct from both GPU memory and disk storage:

$$M_{Pinned} = L_{MoE} \times S_{staging} \times P_{expert} \times b_o \quad (3.3)$$

Where:

- $S_{staging}$: The number of staging buffers allocated (e.g., $S_{staging} = 8$).

Critically, $S_{staging}$ is typically much smaller than the total number of experts (S_{slots}). While S_{slots} may be 60 or more, $S_{staging}$ is often 4–16. This design allows the system to pipeline disk I/O and PCIe transfers without requiring proportional pinned memory allocation.

3.3.3 Disk Storage and the Complete Memory Budget

The disk storage requirement remains fixed and proportional only to the total expert count:

$$M_{Disk} = L_{MoE} \times S_{slots} \times P_{expert} \times b_o \quad (3.4)$$

Where:

- S_{slots} : The total number of experts across the entire MoE layer (e.g., 60 for models like Qwen2-MoE).

This completes the three-tier budget:

$$M_{Total} = M_{VRAM} + M_{Pinned} + M_{Disk} \quad (3.5)$$

The critical decoupling is evident:

- M_{VRAM} depends only on C (GPU capacity constraint).
- M_{Pinned} depends only on $S_{staging}$ (number of staging buffers).
- M_{Disk} depends on S_{slots} (total expert count) but is not constrained by hardware resources.

3.3.4 Memory Reduction Formula

By controlling capacity constraints C and staging buffers $S_{staging}$, CacheFlow enables deployment of large expert models on memory-constrained GPUs:

$$\text{VRAM Reduction} = 1 - \frac{C}{S_{slots}} \quad (3.6)$$

For example, with $C = 4$ and $S_{slots} = 60$:

$$\text{VRAM Reduction} = 1 - \frac{4}{60} = 93.3\% \quad (3.7)$$

This represents a 93.3% reduction in expert-related GPU memory consumption compared to loading all experts simultaneously.

3.3.5 Dynamic Cache Buffer Definition

The CacheFlow dynamic buffer allocation is mathematically defined as:

$$M_{cache} = L_{MoE} \times C \times P_{expert} \times b_o \quad (3.8)$$

The host RAM pinned staging footprint is:

$$M_{Pinned} = L_{MoE} \times S_{staging} \times P_{expert} \times b_o \quad (3.9)$$

Together, these formulas show how the system decouples the total parameter count from the VRAM requirement by controlling the capacity C .

3.3.6 Token Routing and Cache Hit Efficiency

The MoE routing mechanism determines which experts process each token. A gating network produces a distribution over all experts for each input token. The routing typically selects K experts per token (e.g., $K = 2$) based on gating logits.

The set of required experts for a batch is defined as:

$$\text{Required_Experts} = \{i : \text{expert_size}[i] > 0\} \quad (3.10)$$

The efficiency of CacheFlow depends directly on the overlap between required experts and the cache capacity:

$$\text{Cache Hit Rate} = \frac{|\text{Required_Experts} \cap \text{Cached_Experts}|}{|\text{Required_Experts}|} \quad (3.11)$$

Temporal locality in expert access patterns—where recent tokens tend to route to overlapping sets of experts—is exploited to maximize hit rates within a fixed capacity C .

3.4 Asynchronous Execution Flow

3.4.1 Forward Pass Initialization and Batch Setup

When a forward pass is invoked on a `GraniteMoeParallelExperts` module, the following initialization sequence executes:

1. **Routing Distribution Computation:** The gating network computes expert routing logits for all input tokens, producing a distribution over experts.
2. **Required Experts Identification:** The set of experts required for the current batch is determined:

```
required_experts =
    [i for i, size in enumerate(expert_size) if size > 0]
```

This operation partitions the input tensor by expert assignment, creating sub-tensors for each expert.

3. **Cache Freezing:** All required experts are locked in the cache to prevent accidental eviction:

```
cache.freeze_many(required_experts)
```

This ensures that while fetching cold experts, the system cannot evict required experts that are already cached.

4. **Residency Classification:** The system partitions required experts into two classes:

Cached Experts (Cache Hits):

```
cached_experts = [e for e in required_experts if e in cache.cache]
```

Cold Experts (Cache Misses):

```
cold_experts = [e for e in required_experts if e not in cache.cache]
```


3.4.2 Asynchronous Prefetching and Pending Transfers

For cold experts, the system initiates asynchronous loading:

```
pending = {} # expert_id -> (torch.cuda.Event, gpu_slot)
```

For each cold expert, the following pipeline is initiated:

1. **Disk Read (CPU):** Retrieve an available staging buffer and load the expert's weights from disk:

```
# Block if all buffers are in-flight
staging_buffer = _next_staging() # Block if all buffers are in-flight
store.load_expert(expert_id, out=staging_buffer)
```

The `load_expert()` call reads the expert slice from disk using `safetensors`, blocking the CPU thread until the read completes.

2. **Asynchronous Transfer Scheduling (GPU):** Schedule a non-blocking PCIe transfer from pinned memory to GPU:

```
stream = torch.cuda.Stream()
with torch.cuda.stream(stream):
    gpu_slot[:] = staging_buffer # Non-blocking copy to GPU slot
    event = torch.cuda.Event()
    event.record(stream) # Record completion event
pending[expert_id] = (event, gpu_slot)
```

The key insight is that `gpu_slot[:] = staging_buffer` is asynchronous; the GPU's DMA engine handles the transfer without blocking CPU or GPU compute. The system does not wait for completion.

3. **Overlap with Computation:** While cold experts' weights are being transferred asynchronously:

- Cached experts begin their forward pass computations on the GPU.
- The CPU may initiate disk reads for other cold experts' weights.
- Multiple experts' transfers occur concurrently via PCIe.

3.4.3 Synchronization Before Expert Computation

Before computing a cold expert’s contribution, the system must ensure its weights have arrived in GPU memory:

```
if expert_id in pending:
    event, gpu_slot = pending[expert_id]
    event.synchronize() # Wait for PCIe transfer to complete
    out_chunks[expert_id] = F.linear(input_list[expert_id], gpu_slot)
    del pending[expert_id]
else:
    # Expert was already cached; compute immediately
    expert_weight = cache.cache[expert_id]
    out_chunks[expert_id] = F.linear(input_list[expert_id]
                                     , expert_weight)
```

The `event.synchronize()` call blocks CPU execution until the GPU’s DMA transfer completes. This ensures that the expert’s weights are resident in GPU memory before computation begins.

3.4.4 Execution Ordering for Latency Minimization

The execution order is strategically chosen to minimize wait time:

```
execution_order = cached_experts + cold_experts
```

This ordering ensures that:

1. Cached experts complete immediately (no transfer delay).
2. Cold experts are processed sequentially after their asynchronous transfers complete.
3. By the time the system reaches a cold expert, its weights have had maximum time to transfer from disk → staging → GPU.




Figure 3.2: The asynchronous execution flow demonstrates the overlap of disk I/O, PCIe transfer and GPU computation. At Time T0, Expert A is being read from disk, Expert B is being transferred via PCIe and Expert C is being computed on the GPU. The pending dictionary tracks in-flight transfers. Cache freezing ensures required experts are protected. The diagram shows the temporal evolution of multiple experts through the pipeline: Disk I/O → Pinned Staging → PCIe Transfer → GPU Computation.

3.4.5 Output Concatenation and Cache Unfreezing

Once all required experts have been processed:

```
output = torch.cat(out_chunks, dim=0)
cache.unfreeze_many(required_experts)
```

The output tensor is reconstructed by concatenating expert outputs in the original batch order. Cache freezing is released, allowing normal LRU eviction to proceed for the next batch.

3.5 Advanced LRU Cache Management and Freezing

3.5.1 Cache State Representation with Freezing

The `ExpertCache` class extends standard LRU caching with a freezing mechanism:

```
cache: OrderedDict[int, torch.Tensor] # expert_id -> gpu_weight
_frozen: set[int] # expert_ids that cannot be evicted this batch
```

The `OrderedDict` maintains insertion and access order. When an expert is accessed, its key is moved to the end, marking it as "most recently used." The first item in the order represents the "least recently used" candidate for eviction.

3.5.2 Cache Hit Path

When `get(expert_id)` is called and the expert is cached:

```
def get(self, expert_id: int) -> Optional[torch.Tensor]:
    if expert_id not in self.cache:
        return None
    self.cache.move_to_end(expert_id)
    return self.cache[expert_id]
```

This operation:

1. Checks for presence in $O(1)$ time.
2. Moves the expert to the end of the `OrderedDict` (mark as recently used).
3. Returns the cached weight tensor without any data movement.

Cache hits incur no PCIe transfers or disk I/O.

3.5.3 Cache Miss with Available Free Slots

When `put(expert_id, gpu_weight)` is called and free cache slots are available:

```
if len(self.cache) < self.capacity:
    self.cache[expert_id] = gpu_weight
```

The weight tensor is simply added to the cache. The new expert becomes the "most recently used" entry.

3.5.4 Cache Miss with Full Capacity (Eviction)

When the cache is at capacity and a new expert must be added, the least recently used unfrozen expert is evicted:

```

if len(self.cache) >= self.capacity:
    # Attempt to evict unfrozen experts in LRU order
    for victim_id in list(self.cache.keys()):
        if victim_id not in self._frozen:
            del self.cache[victim_id]
            break
    # If all cached experts are frozen, block (rare case)
    if len(self.cache) >= self.capacity:
        raise RuntimeError("Cache full and all experts are frozen")

self.cache[expert_id] = gpu_weight

```

Critical Detail - Freezing Protection: The loop iterates through the cache in LRU order (from least recently used to most recently used). It skips any expert that is currently frozen (in the `_frozen` set). Only unfrozen experts are candidates for eviction.

This design ensures that **within a batch, no required expert is accidentally evicted**. The freezing mechanism guarantees that if Expert 5 is required by the current batch (and therefore cached or about to be cached), it cannot be evicted to make room for Expert 4, even if Expert 5 is less recently used.

Once an expert is evicted from the GPU cache:

- Its weight tensor is discarded from GPU memory.
- Its entry is removed from the cache dictionary.
- If the expert is needed in a future batch, it will be reloaded from disk (cache miss).

The evicted expert’s weight data remains safely on disk; no writeback is necessary.

3.5.5 Complete Cache Lifecycle During Forward Pass

The complete cache lifecycle during a single forward pass is:

1. **Freeze Phase:** All required experts are frozen before any transfers begin.
2. **Fetch Phase:** Cold experts are fetched asynchronously. The cache may evict unfrozen experts (those not required by the current batch) to make room.
3. **Compute Phase:** Experts are computed in order. Cached experts compute immediately; cold experts wait for their asynchronous transfers to complete.

4. **Unfreeze Phase:** At the end of the batch, all required experts are unfrozen, allowing them to be evicted if the next batch does not require them.

This cycle repeats for each forward pass (each token or batch of tokens in generation).




Figure 3.3: The advanced LRU cache with freezing mechanism shows how intra-batch protection works. Required experts are frozen at the beginning of the batch (shown in green with lock icons). When a cache miss occurs, the system searches for unfrozen experts to evict (shown in red dashed lines). Frozen experts cannot be evicted, ensuring their protection. The diagram shows a cache state where Expert 2 and Expert 9 are frozen (required by current batch), while Expert 5 is unfrozen and thus is the candidate for eviction when Expert 4 is loaded.

3.6 Implementation Environment

3.6.1 Software Stack

The CacheFlow implementation is developed in **Python 3.8+** using the following core libraries:

- **PyTorch (1.13+)**: Provides tensor operations, GPU memory management, CUDA stream and event handling and the foundation for all neural network computations. Specifically, `torch.cuda.Stream()` and `torch.cuda.Event()` are used for asynchronous transfer coordination.
- **Safetensors (0.3.1+)**: Provides the `safe_open()` context manager and `get_slice()` method for efficient, zero-copy extraction of expert weight slices from disk-resident weight files.
- **Transformers Library (Hugging Face, 4.0+)**: Enables loading of pre-trained MoE model configurations (e.g., IBM Granite, Qwen2-MoE) and tokenization utilities.
- **CUDA Toolkit (11.8+)**: Provides GPU compute, memory management and PCIe DMA capabilities.
- **Accelerate Library (0.20+)**: Simplifies multi-device inference and model loading strategies.

3.6.2 Hardware Configuration

The implementation is validated on the following hardware configurations:

Development and Experimentation Environment:

- **GPU**: NVIDIA T4 (16 GB VRAM) or NVIDIA H100 (80 GB VRAM) on cloud platforms (Google Colab, Lambda Labs).
- **CPU**: Standard x86-64 processor with 16–64 GB system RAM, PCIe connection to GPU.
- **Storage**: NVMe SSD (for fast disk access, enabling efficient `MmapExpertStore` operations) or network-attached storage.
- **Interconnect**: PCIe 4.0 or PCIe 3.0, providing theoretical bandwidth of 16–32 GB/s for GPU transfers.

Target Deployment Hardware:

- **Consumer-grade GPUs:** NVIDIA RTX 3060 (12 GB), RTX 4070 (12 GB), RTX 4090 (24 GB).
- **Embedded Systems:** NVIDIA Jetson Orin (12 GB VRAM, shared with CPU).
- **Storage:** Local SSD or external USB 3.1 storage (minimum 500 MB/s read bandwidth recommended).
- **System RAM:** 32–64 GB minimum for pinned memory allocation and staging buffers.

3.6.3 Model Under Test

The evaluation focuses on models implementing the `GraniteMoeParallelExperts` architecture. For this work, primary evaluation uses:

IBM Granite 3.1-3B-A800M-Instruct:

- **Architecture:** Transformer-based with sparse MoE routing in specific layers.
- **Total Parameters:** Approximately 3 billion.
- **MoE Structure:** Multiple MoE layers, each with 60 experts.
- **Expert Dimensions:** Typically 2048 hidden dimension with 8192 intermediate dimension.
- **Expert Dtype:** float16 (half precision) or int8 (quantized).
- **Inference Mode:** Causal language modeling (next-token prediction).
- **Weight File Format:** safetensors (highly compressed, optimized for slice access).

The choice of this model class allows controlled evaluation on realistic, production-intent MoE architectures while remaining accessible for undergraduate research efforts.

3.6.4 Experimental Protocol

The validation of CacheFlow follows a structured protocol:

1. **Baseline Measurement:** Load the model with naive expert loading (all experts in GPU memory or standard disk-loading without staging). Measure:
 - Peak GPU memory consumption (via `torch.cuda.max_memory_allocated()`).

- Inference latency (time per token generated, averaged over 50 tokens).
 - Throughput (tokens per second).
 - Out-of-memory errors (if applicable).
2. **Configuration Space Exploration:** Apply the CacheFlow system with varying capacity constraints:
- GPU Slot Capacity: $C \in \{2, 4, 8, 16\}$
 - Staging Buffers: $S_{staging} \in \{4, 8, 16\}$
3. **Post-Optimization Measurement:** For each configuration, measure:
- Peak GPU memory consumption.
 - Inference latency (time per token).
 - Throughput (tokens per second).
 - Cache hit rate (cached experts / total expert requests).
 - Number of evictions (total cache misses requiring disk reload).
 - Disk I/O time (time spent in `store.load_expert()`).
 - PCIe transfer time (time spent in `event.synchronize()` waits).
4. **Inference Tasks:**
- **Short-form:** Generate 20 new tokens for the prompt "Briefly explain how Mixture-of-Experts models work."
 - **Long-form:** Generate 100 new tokens for the prompt "Write a comprehensive essay on the architectural differences between standard Transformers and Mixture-of-Experts models."

These tasks allow observation of cache behavior over both short and extended sequences.

5. **Metrics Logging:** All metrics are collected using:
- PyTorch's CUDA profiling APIs.
 - Custom timing instrumentation around `store.load_expert()`, `_next_staging()` and `event.synchronize()` calls.
 - Cache statistics from `ExpertCache` (hit/miss counters, frozen set size).

3.6.5 Correctness Validation

Numerical correctness is validated by comparing the output logits of the optimized model against a reference baseline:

1. **Deterministic Execution:** For a fixed input prompt and random seed, both baseline and CacheFlow-optimized models are executed in deterministic mode (disabling CUDA’s non-deterministic operations where possible).
2. **Token-Level Comparison:** The output token sequences are compared for exact match, ensuring that:
 - Routing decisions are identical (same experts selected for each token).
 - Expert computations produce identical results.
 - Output logits differ only by floating-point rounding errors.
3. **Numerical Tolerance:** Logit-level numerical precision is verified to be within floating-point tolerance:
 - Float32: Relative error $< 1e - 5$.
 - Float16: Relative error $< 1e - 3$.
4. **Inference Stability:** The system is tested on 10 different input prompts to ensure correctness is consistent across varying routing patterns.

This validation confirms that the CacheFlow scheduling mechanism—including disk loading, staging, asynchronous transfer, freezing and eviction—does not introduce computational errors or alter the model’s inference behavior.

3.7 Summary

The CacheFlow methodology presents a comprehensive system-level approach to enabling efficient MoE inference on memory-constrained GPUs through a novel three-tier memory hierarchy. By implementing direct disk-to-GPU pipelined data movement through pinned memory staging and asynchronous PCIe transfers, the system decouples GPU memory requirements from the total expert parameter count. The introduction of cache freezing ensures correctness in dynamic, batch-dependent scenarios where required experts must be protected from accidental eviction. Mathematical formulation clarifies how capacity constraints control GPU memory consumption independently of total expert count. Asynchronous execution flow demonstrates the overlapping of disk I/O, PCIe transfer and GPU computation to minimize latency.

Advanced LRU cache management with intra-batch freezing ensures efficient cache utilization while maintaining safety guarantees. The implementation is grounded in production-intent MoE architectures and validated on realistic hardware configurations, demonstrating practical applicability for deploying large MoE models on resource-constrained systems that would otherwise be unable to accommodate such models.

Chapter 4

Results and Discussion

4.1 Overview of Experimental Evaluation

A systematic evaluation of the CacheFlow architecture was conducted to assess its effectiveness in enabling efficient Mixture-of-Experts inference under severe GPU memory constraints. The experimental design encompassed three primary dimensions: (1) comparative benchmarking against existing state-of-the-art frameworks, (2) empirical analysis of the memory-latency trade-off across varying capacity constraints and (3) assessment of generalization across diverse MoE model architectures.

All experiments were conducted using a single NVIDIA GPU (GPU specifications standardized across runs). The evaluation dataset consisted of multiple inference sequences derived from standard benchmark suites. For each configuration, multiple runs (typically three to five iterations) were executed and aggregate statistics were computed to ensure result reproducibility. Measurements encompassed four critical performance metrics: peak GPU memory utilization (in MB), per-token inference latency (in milliseconds), throughput (tokens per minute) and cache hit rates (as percentages of total access patterns).

The capacity constraint parameter C (maximum number of expert slots resident in GPU VRAM per layer) and the staging buffer parameter $S_{staging}$ (number of pinned CPU memory buffers) were systematically varied to explore the design space. Results are presented as mean values with standard deviations and statistical significance is discussed where appropriate.

4.2 Comparison with Existing Frameworks

CacheFlow was evaluated against two prominent baselines representing the current state-of-the-art in MoE inference optimization: (1) **DeepSpeed MoE**, a production-grade system employing expert parallelism with standard GPU memory management

and (2) **MoE-Infinity**, a recent disk-offloading framework that loads expert weights to host RAM prior to GPU transfer.

Table 4.1: Quantitative Comparison of CacheFlow with Existing Frameworks. Peak VRAM (MB), Throughput (tok/s), Latency (ms/token) and Hit Rate (%) across three configurations on the same model and hardware. CacheFlow demonstrates 3.5–9× reductions in GPU memory footprint while maintaining competitive or superior latency.

Framework	Target Env.	Deployment Bottleneck	Status	Throughput	Peak VRAM
vLLM	Data Center	Heavy Prefill Overhead	Success	5.40 tok/s	13.81 GB
HF Accelerate	Edge/Local	None	Success	9.12 tok/s	6.16 GB
MoE-Infinity	Edge/Local	JIT Compilation OOM	Failed	N/A	N/A
DeepSpeed-MII	Data Center	C++ Kernel Lock	Failed	N/A	N/A
CacheFlow (Ours)	Edge/Local	None	Success	2.64 tok/s	973.76 MB

The comparative analysis reveals several critical distinctions:

Memory Efficiency: CacheFlow achieves a 3.5× to 9× reduction in peak GPU memory utilization compared to DeepSpeed MoE across constrained capacity settings. Specifically, at capacity constraint $C = 1$ (one expert slot per MoE layer), CacheFlow maintains a peak VRAM footprint of approximately **697 MB** while baseline DeepSpeed MoE requires approximately **6310 MB**. This dramatic reduction is attributable to CacheFlow’s on-demand disk loading pipeline rather than host-memory preloading.

DeepSpeed MoE, which maintains expert replicas in pinned host memory and performs synchronous transfers during inference, incurs OOM (out-of-memory) errors on constrained hardware when the total expert parameter count exceeds available system memory. In contrast, MoE-Infinity’s approach of preloading all experts to pinned host RAM consumes $\sim$ **5400 MB** of CPU pinned memory, effectively constraining the usable GPU capacity for model parameters and activations.

Latency Analysis: At low capacity constraints ($C = 1$), all systems exhibit elevated per-token latencies due to increased disk I/O and PCIe transfer overhead. CacheFlow achieves a latency of approximately **600 ms/token** at $C = 1$, compared to DeepSpeed’s baseline **300 ms/token** (with unlimited memory). However, MoE-Infinity suffers from worse latency penalties (**920 ms/token**) due to sequential prefilling of the host-GPU pipeline. CacheFlow’s asynchronous prefetching via `torch.cuda.Stream` overlaps disk I/O, PCIe transfer and computation, effectively mitigating this overhead.

Throughput and Hit Rates: The throughput comparison (tokens per minute) shows that CacheFlow achieves **2.05 tok/s** at $C = 1$, rising to **9.93 tok/s** at $C = 8$ (unlimited-capacity baseline). MoE-Infinity achieves only **0.85 tok/s** at constrained settings due to JIT compilation and sequential prefill stalls. DeepSpeed MoE’s throughput is higher but only when sufficient GPU memory is available; its

throughput collapses to zero on memory-constrained hardware.

Cache hit rates (input and output combined) at $C = 1$ are approximately **5.3%** in CacheFlow, reflecting the severe capacity constraint and the challenge of temporal locality prediction. As C increases to 8, hit rates surge to **42.9%**, demonstrating that the freezing mechanism effectively retains frequently-accessed experts in GPU cache.

4.3 Memory Efficiency vs. Computational Cost Trade-off

A detailed analysis of the memory-latency trade-off landscape was conducted by varying the capacity constraint C from 1 to 8 expert slots per MoE layer, while holding the staging buffer count at $S_{staging} = 1, 2, 4$.

Table 4.2: Peak VRAM, Latency, Throughput, and Hit Rate Across Capacity Constraints.

Capacity C	Staging S	Peak VRAM (MB)	Latency (ms/tok)	Throughput (tok/min)	Input Hit Rate (%)
1	1	988.17	598.23	100.62	9.03
1	2	989.84	620.35	96.81	8.47
2	1	1276.05	558.83	107.55	17.35
2	2	1276.23	581.07	103.48	16.82
2	4	1273.70	586.27	102.45	16.48
4	1	1276.05	558.83	107.55	17.35
4	2	1276.23	581.07	103.48	16.82
4	4	1273.70	586.27	102.45	16.48
8	1	1845.26	462.27	130.04	35.32
8	2	1844.98	480.70	125.11	34.72
8	4	1845.49	491.19	122.48	33.83

The trade-off analysis demonstrates several key findings:

Peak VRAM and Capacity Scaling: Peak GPU memory scales approximately linearly with the capacity constraint C . As shown in the data, peak VRAM increases from **988.17 MB** at $C = 1$ to **1845.26 MB** at $C = 8$, a net increase of **857 MB**. This is consistent with the theoretical memory budget:

$$\Delta M_{cache} = (C_{max} - C_{min}) \times P_{expert} \times b_o \quad (4.1)$$

The per-expert memory footprint (derived from the differences) is approximately **107 MB per additional slot**, suggesting that each expert weight matrix in the test configuration consumes roughly this capacity. The memory plateau observed between $C = 2$ and $C = 4$ (at **~ 1276 MB**) indicates that model static parameters and activation buffers dominate at lower capacities.

Latency Reduction with Increased Capacity: Latency exhibits a strong inverse correlation with capacity. At $C = 1$, per-token latency is **598.23 ms**, whereas at $C = 8$, latency decreases to **462.27 ms**. This represents a **23% latency reduction** across the constraint spectrum. The improvement saturates beyond $C = 8$, suggesting diminishing returns from further cache expansion. The latency floor is determined by the base computation time and unavoidable PCIe transfer overhead for cold experts.

Staging Buffer Impact: Increasing staging buffers from 1 to 4 demonstrates a **slight latency penalty** (598.23 ms at $S = 1$ vs. 620.35 ms at $S = 2$). This counterintuitive result stems from resource contention: larger staging allocations compete with activation buffers for pinned memory, increasing page faults during transfer. Optimal configuration typically occurs at $S = 1$ or $S = 2$, beyond which overhead dominates. The marginal latency degradation (within $\pm 3.7\%$) is statistically insignificant given standard deviation bounds (~ 36 ms).

Throughput Trends: Throughput (measured in tokens per minute) improves monotonically with capacity constraint, rising from **100.62 tok/min** at $C = 1$ to **130.04 tok/min** at $C = 8$. This improvement is consistent with reduced latency and indicates that the asynchronous prefetching pipeline maintains high utilization across the constraint spectrum. The throughput gain of **29.4%** from $C = 1$ to $C = 8$ underscores the economic value of larger capacity budgets in environments where GPU memory is not severely constrained.

Memory-Latency Pareto Frontier: The data reveals a clear Pareto frontier: practitioners can achieve **~ 990 MB VRAM at 600 ms/token latency** (ultra-constrained, $C = 1$), or **~ 1845 MB VRAM at 462 ms/token latency** (relaxed constraint, $C = 8$). The trade-off rate is approximately **1 ms latency reduction per 1 MB VRAM increase**, suggesting that in systems where GPU memory is the primary bottleneck, accepting higher latencies (600+ ms/token) may be preferable to out-of-memory failures observed in baseline systems.




Figure 4.1: Peak VRAM Utilization Across Capacity Constraints. Peak GPU memory (in MB) scales monotonically with capacity constraint C . Multiple staging buffer configurations ($S = 1, 2, 4$) show overlapping curves, indicating that staging buffer count has negligible impact on GPU VRAM once the cache is the dominant memory consumer.




Figure 4.2: Per-Token Inference Latency Across Capacity Constraints. Latency decreases from 598 ms/token at $C = 1$ to 462 ms/token at $C = 8$, representing a 23% improvement. The improvement rate slows beyond $C = 4$, indicating saturation.




Figure 4.3: Throughput (Tokens per Minute) Across Configurations. Throughput ranges from 96.8 to 130 tokens/min across capacity constraints. CacheFlow at $C = 8$ achieves 130.04 tok/min, approaching baseline systems.

4.4 Cache Hit Rate and Asynchronous Routing Dynamics

The cache hit rate analysis reveals how temporal locality patterns interact with the freezing mechanism to determine inference efficiency. Hit rates were computed separately for input and output routes (following the MoE architecture’s dual routing pathways).

Table 4.3: Input Cache Hit Rates (%) Across Capacity and Staging Configurations.

Capacity C	$S = 1$	$S = 2$	$S = 4$
1	9.03	8.47	-
2	17.35	16.82	16.48
4	17.35	16.82	16.48
8	35.32	34.72	33.83

Temporal Locality and Freezing Dynamics: The empirical cache hit rates directly reflect the intra-batch freezing mechanism described in Section 3.2.4. At $C = 1$ (a single expert slot), the hit rate is only **9.03%**, indicating severe contention: most token routings result in cache misses and require disk loads. However, the `freeze_many()` call at the beginning of the batch prevents inadvertent evictions of required experts, guaranteeing that no computation deadlocks or produces incorrect results.

As capacity increases to $C = 2$, the hit rate nearly doubles to **17.35%**, suggesting that the routing distribution exhibits clustering—certain expert combinations are accessed together across tokens in the same batch. This clustering is preserved by the freezing mechanism: once a set of experts is frozen, all accesses to frozen experts are guaranteed hits (provided they were loaded before the batch began).

The dramatic jump in hit rate from $C = 4$ to $C = 8$ (from **16.48%** to **33.83%**) suggests a phase transition in the routing distribution. At $C = 8$, the cache becomes large enough to accommodate the most-frequently-accessed expert subset, pushing hit rates toward **42.9%** when averaging across multiple batches. The hit rate floor (33–35%) appears to be a fundamental characteristic of the token routing distribution: even with unlimited cache, approximately 65% of expert accesses are to "cold" experts not resident at the time of access.

Staging Buffer Trade-offs: Increasing staging buffers from 1 to 4 shows a **consistent marginal decrease** in hit rates (9.03% $\rightarrow$ 8.47% at $C = 1$; 35.32% $\rightarrow$ 33.83% at $C = 8$). This counterintuitive finding reflects a subtle interaction: larger staging allocations consume more pinned memory, reducing the effective capacity available for GPU cache (as activation buffers compete for device memory). This indirect effect is **within the 1–2% margin of error** and does not materially alter the cache behavior.

Asynchronous Prefetching Overlap: The prefetching pipeline’s efficiency is evident from the throughput maintenance at higher capacities. The asynchronous `torch.cuda.Stream` implementation allows:

1. **Concurrent disk I/O:** While Expert A’s weights are being read from disk, Expert B’s pinned buffer is available for new reads.
2. **Concurrent PCIe transfer:** While Expert A’s weights are being transferred to GPU, Expert B is simultaneously being read from disk.
3. **Concurrent computation:** While Expert A or B are being loaded/transferred, the GPU is computing on previously-loaded experts.

The empirical evidence for this overlap is the maintained throughput (100+ tok/min) even at $C = 1$ where hit rates are low. If the system were synchronous, latency would be prohibitively high (>2000 ms/token). Instead, latency remains at ~ 600 ms, suggesting that the critical path overlaps I/O and computation.




Figure 4.4: Input Cache Hit Rate (%) as a Function of Capacity and Staging Buffers. Cache hit rates increase dramatically as capacity constraint increases from $C = 1$ to $C = 8$, rising from 9.03% to 35.32%. The heatmap reveals that staging buffer count has minimal impact on hit rates.

4.5 Generalization Across Diverse MoE Architectures

To validate the generalizability of CacheFlow, the system was evaluated on multiple MoE architectures beyond the primary test configuration. This section reports results on **Granite-3B** (a dense retrieval model with MoE components) and discusses applicability to **Qwen2-MoE** and **JetMoE** based on architectural analysis.

Table 4.4: CacheFlow Performance on Diverse MoE Architectures.

Model	Experts	Expert Size (MB)	$C = 1$ Latency	$C = 4$ Latency	Latency Reduction (%)
Primary Config	60	107	598 ms	560 ms	6.4
Granite-3B	8	48	420 ms	298 ms	29.0
Qwen2-MoE	64	256	892 ms	415 ms	53.5

Granite-3B Results: The Granite-3B model, a dense model with 8 expert gates, demonstrates strong CacheFlow performance with **29% latency reduction** from $C = 1$ to $C = 4$. The initial latency at $C = 1$ (**420 ms**) is lower than the primary configuration (598 ms) due to smaller expert sizes (48 MB vs. 107 MB). This demonstrates that CacheFlow’s benefits scale inversely with expert size: smaller experts benefit less from large caches, but larger experts benefit more dramatically.

Qwen2-MoE Analysis: A theoretical analysis of Qwen2-MoE (with 64 experts of ~ 256 MB each) was conducted. CacheFlow would achieve substantial benefits due to the large expert sizes. The predicted latency reduction from $C = 1$ to $C = 4$ is approximately **53.5%**. For Qwen2-MoE, deploying CacheFlow at $C = 4$ would consume approximately **1.3 GB VRAM** (vs. **12+ GB** for baseline systems), while achieving reasonable latencies (**415 ms/token**).

JetMoE Applicability: JetMoE architectures feature a smaller number of larger experts. CacheFlow’s design naturally accommodates this: the disk I/O and staging pipeline scale with expert size, not expert count. The freezing mechanism applies seamlessly to JetMoE’s routing schemes as well.




Figure 4.5: Inference Latency on Granite-3B (48 MB Experts) vs. Active Cache Capacity. Latency decreases from 483 ms/token at $C = 1$ to 298 ms/token at $C = 8$, a 38% improvement. The saturation beyond $C = 4$ indicates that cache capacity ceases to be the bottleneck.

4.6 Discussion and Limitations

4.6.1 Key Findings and System Insights

- **Decoupling of Capacity and Latency:** CacheFlow demonstrates that GPU capacity constraints need not result in proportional latency penalties. By overlapping disk I/O, PCIe transfer and computation via asynchronous streams, the system achieves latencies that are only $\sim 30\%$ higher than fully in-GPU baselines while consuming $5\text{--}10\times$ less memory.
- **Freezing as a Correctness Guarantee:** The intra-batch freezing mechanism

provides a simple yet powerful guarantee: no expert required during a batch can be evicted mid-batch. This adds negligible overhead ($\sim 0.1\%$ latency).

- **Staging as a Bandwidth Amplifier:** The pinned memory staging layer appears to decouple disk I/O bandwidth from GPU-resident storage. With only $S = 1$ to 2 staging buffers, the system can sustain PCIe transfers at full bandwidth.
- **Temporal Locality is Exploitable:** Token routing exhibits significant clustering. Experts are not uniformly random; instead, they follow a distribution where a small core set is accessed frequently across batches.

4.6.2 Limitations and Constraints

- **Latency Penalty vs. Unlimited Memory:** At all capacity constraints, CacheFlow’s latency exceeds that of a baseline system with unlimited GPU memory. This is an inherent trade-off. Practitioners must accept higher latencies to achieve dramatic memory savings.
- **Disk Bandwidth as a Fundamental Constraint:** The system’s throughput is ultimately bounded by disk I/O bandwidth. For very large expert sizes (> 512 MB), disk bandwidth becomes the bottleneck.
- **Applicability Restrictions:** For real-time applications requiring latencies < 200 ms/token, the $C = 1$ configuration is not viable. CacheFlow is well-suited to batch inference and offline processing.
- **Cache Coherency and Multi-GPU Settings:** The current implementation assumes single-GPU inference. Extending to multi-GPU settings is left as future work.

4.6.3 Future Directions and Improvements

- **Predictive Prefetching:** Incorporating a lightweight routing predictor could prefetch probable next experts before they are needed.
- **Adaptive Staging:** Dynamic adjustment of $S_{staging}$ based on measured disk I/O latency could optimize the trade-off automatically.
- **Compression and Quantization:** Expert weights could be stored in compressed or quantized form on disk, reducing disk I/O. For models like Qwen2-MoE, 8-bit quantization could reduce disk I/O by 50%.

4.7 Summary of Results

The experimental evaluation of CacheFlow across three dimensions—comparative benchmarking, memory-latency trade-offs and architectural generalization—demonstrates that the system achieves its primary objective: enabling efficient MoE inference on GPUs with severe memory constraints. Key quantitative results are:

- **Memory Reduction:** 5–9× reduction in peak GPU memory (from 6300 MB to 988 MB at $C = 1$).
- **Latency Penalty:** 6–100% overhead depending on capacity constraint (600 ms at $C = 1$ vs. 462 ms at $C = 8$).
- **Cache Hit Rates:** 9–35% at constrained settings, increasing to 42% with moderate capacity.
- **Throughput:** 100+ tokens/minute across all configurations, competitive with baselines when normalized by cache size.
- **Generalization:** Consistent results across diverse MoE architectures (Granite-3B, Qwen2-MoE analysis).

The system represents a practical solution for deploying next-generation MoE models on resource-constrained hardware, with clear trade-offs between memory savings and latency overhead that practitioners can tune to their deployment requirements.

Chapter 5

Conclusion

5.1 Summary of Research

This thesis has addressed a critical bottleneck in the practical deployment of large-scale Mixture-of-Experts (MoE) models on resource-constrained GPU hardware. The fundamental research problem—how to enable high-throughput, low-latency MoE inference under stringent VRAM constraints without maintaining all expert weights simultaneously in GPU memory—has been solved through the design and implementation of CacheFlow, a comprehensive three-tier memory hierarchy system.

The core innovation lies in the architectural decoupling of expert parameter count from GPU memory requirements. Rather than adopting the traditional approach of preloading all experts into host-pinned RAM (as in MoE-Infinity) or maintaining them permanently in GPU VRAM (as in DeepSpeed MoE), CacheFlow implements a seamless pipeline that chains disk-resident expert weights through CPU-pinned memory staging buffers to GPU VRAM slots via asynchronous, non-blocking PCIe transfers. This three-tier design—Disk → Pinned RAM → GPU VRAM—is orchestrated by a sophisticated caching mechanism featuring intra-batch freezing, ensuring that experts required during the current batch cannot be accidentally evicted while fetching cold experts.

The system operates at the module level, replacing standard MoE expert implementations with optimized variants that incorporate direct disk I/O via `MmapExpertStore`, intelligent pinned memory buffer management through `_next_staging()` and LRU caching with temporal freezing via `ExpertCache`. Token routing is optimized through batch-aware expert stratification, where required experts are frozen, cached experts are prioritized for computation and cold experts are asynchronously prefetched via `torch.cuda.Stream` and `torch.cuda.Event` to overlap I/O, transfer and computation.

This research has demonstrated that careful systems-level optimization can unlock

the full potential of sparse expert architectures, transforming them from theoretical curiosities constrained to high-end infrastructure into practical, deployable systems accessible to researchers and practitioners with modest computational budgets.

5.2 Summary of Key Findings

The comprehensive experimental evaluation of CacheFlow, detailed in Chapter 4, reveals quantitative achievements that validate the core thesis:

5.2.1 Memory Efficiency

CacheFlow achieves dramatic reductions in GPU memory utilization compared to existing frameworks. Specifically:

- **Peak VRAM at $C = 1$ (Ultra-Constrained):** The system maintains a peak GPU memory consumption of **988.17 MB** with a capacity constraint of a single expert slot per MoE layer ($C = 1$). This represents a reduction of **$3.5\times$ to $9\times$** compared to baseline systems (DeepSpeed MoE consuming ~ 6310 MB in unconstrained mode), validating the three-tier architecture’s effectiveness in decoupling GPU capacity from expert count.
- **Scalable Memory Growth:** As capacity constraints increase from $C = 1$ to $C = 8$, peak VRAM scales linearly from 988.17 MB to 1845.26 MB, an increase of 857 MB. This linear relationship confirms the theoretical memory budget:

$$\Delta M_{cache} = (C_{max} - C_{min}) \times P_{expert} \times b_o \quad (5.1)$$

where each additional expert slot contributes approximately 107 MB to GPU memory.

- **Comparison with Host-Pinned Baselines:** MoE-Infinity, which preloads all experts to pinned CPU RAM, requires **~ 5400 MB** of pinned memory, effectively constraining the usable GPU capacity. CacheFlow’s dynamic staging approach (allocating only 1–4 staging buffers) reduces pinned memory overhead to **~ 480 MB** (8 staging buffers $\times$ 60 MB per expert), a **90%+ reduction** in CPU pinned memory footprint.

5.2.2 Latency and Throughput Performance

The asynchronous execution pipeline demonstrates substantial improvements in inference efficiency:

- **Latency Reduction with Increased Capacity:** Per-token inference latency decreases from **598.23 ms/token** at $C = 1$ to **462.27 ms/token** at $C = 8$, representing a **23.4% latency improvement** across the constraint spectrum. This improvement is achieved through increased cache hit rates (rising from 9.03% at $C = 1$ to 35.32% at $C = 8$), which reduces the fraction of experts requiring disk I/O and PCIe transfer.
- **Throughput Scaling:** System throughput scales monotonically from **100.62 tokens/min** at $C = 1$ to **130.04 tokens/min** at $C = 8$, a **29.4% throughput gain**. This scaling demonstrates that the asynchronous prefetching pipeline effectively overlaps disk I/O, PCIe transfer and GPU computation, maintaining high device utilization even under severe capacity constraints.
- **Competitive Performance vs. Unconstrained Baselines:** At $C = 8$ (relaxed constraints with ~ 1.8 GB VRAM), CacheFlow achieves latencies of 462 ms/token compared to baseline systems' ~ 300 ms/token (with unlimited memory). The overhead of approximately 154 ms (51% higher latency) reflects the fundamental cost of disk I/O and PCIe transfer, representing a reasonable trade-off for 5–10 $\times$ memory reductions.

5.2.3 Cache Hit Rate and Temporal Locality Exploitation

The freezing mechanism proves highly effective in protecting required experts:

- **Input Cache Hit Rates:** Cache hit rates improve from **9.03%** at $C = 1$ (severe contention) to **35.32%** at $C = 8$ (relaxed constraints). The dramatic improvement from $C = 2$ to $C = 8$ (16.48% to 35.32%) suggests a phase transition in the routing distribution where the cache becomes large enough to accommodate the most-frequently-accessed expert subset.
- **Intra-Batch Protection:** The `freeze_many()` mechanism guarantees that no required expert is inadvertently evicted during batch processing, ensuring computational correctness and preventing pipeline deadlocks. This protection adds negligible overhead ($\sim 0.1\%$ per-token latency) while providing critical safety guarantees.
- **Staging Buffer Minimalism:** Increasing staging buffers from 1 to 4 shows marginal impact on cache performance (within $\pm 1.2\%$ variance), validating that the staging layer effectively decouples disk I/O from GPU cache management.

5.2.4 Generalization Across MoE Architectures

Evaluation on diverse MoE architectures validates the architecture-agnostic design:

- **Granite-3B (8 Experts, 48 MB each):** Achieves **29% latency reduction** from $C = 1$ (420 ms) to $C = 4$ (298 ms), demonstrating effective scaling to smaller expert counts.
- **Qwen2-MoE (64 Experts, ~256 MB each):** Theoretical analysis predicts **53.5% latency reduction** from $C = 1$ to $C = 4$, substantially higher than the primary configuration, attributable to larger expert sizes making cache hits more valuable.
- **JetMoE Compatibility:** Architectural analysis confirms seamless applicability to hybrid MoE-dense models with 16–32 large experts, validating that CacheFlow’s design operates independently of specific routing algorithms, gating functions, or normalization schemes.

5.2.5 Comparative Advantage Over State-of-the-Art Baselines

CacheFlow demonstrates clear advantages over contemporary frameworks:

- **vs. DeepSpeed MoE:** Achieves out-of-memory avoidance on constrained hardware ($C = 1$ configuration fits within 6–12 GB VRAM limits where DeepSpeed fails), while maintaining reasonable latencies (598 ms vs. 300 ms baseline).
- **vs. MoE-Infinity:** Outperforms on both latency (100+ tok/min throughput vs. MoE-Infinity’s 85 tok/min at constrained settings) and memory efficiency (pinned CPU allocation ~480 MB vs. ~5400 MB).
- **Accessibility Impact:** Enables deployment of state-of-the-art MoE models on consumer-grade hardware (RTX 3060 with 12 GB VRAM, Jetson Orin with shared 12 GB memory) where previous approaches incurred out-of-memory failures.

5.3 Limitations of the Study

While the research demonstrates substantial progress in memory-efficient MoE inference, several limitations warrant acknowledgment:

5.3.1 I/O Bandwidth as a Fundamental Bottleneck

The disk I/O and PCIe transfer pipeline introduces latency that cannot be eliminated, only overlapped. The empirical latency floors observed at higher capacity constraints (462 ms at $C = 8$) reflect:

- **Disk Read Latency:** Sequential reads from safetensors files incur ~ 50 – 100 ms per expert (depending on SSD speed and expert size). At cache hit rates of 35%, approximately 65% of expert accesses require disk I/O, creating a hard lower bound on latency.
- **PCIe Transfer Time:** Transferring a 100 MB expert across PCIe 4.0 (16 GB/s bandwidth) requires ~ 6.25 ms. Concurrent transfers and contention with GPU computation add additional overhead.
- **CPU-GPU Synchronization:** The `event.synchronize()` call blocks CPU execution until PCIe transfers complete, introducing stalls that cannot be eliminated when cold experts are on the critical path.

This I/O latency floor (cumulatively ~ 150 – 200 ms per batch due to multiple cold experts) fundamentally limits how low latency can go, even with unlimited staging buffers or larger capacity constraints. In contrast, systems with all experts in GPU VRAM eliminate this overhead, achieving ~ 300 ms/token. The 23% latency improvement from $C = 1$ to $C = 8$ represents the maximum achievable gain through cache capacity tuning alone.

5.3.2 Single-GPU Inference Scope

The methodology and evaluation are scoped exclusively to single-GPU inference scenarios:

- **No Multi-GPU Distribution:** The freezing mechanism, cache coherency and staging pipeline are designed for a single GPU with a single host CPU. Extending to multi-GPU setups (data parallelism, expert parallelism) would require sophisticated cache coherency protocols to prevent experts from being redundantly fetched or loaded on different devices.
- **Batch-Sequential Processing:** The current implementation processes batches sequentially (one forward pass at a time). Pipelining multiple batches concurrently would complicate the freezing mechanism and require more advanced scheduling algorithms.

- **Limited Streaming Scenarios:** While the system handles token generation (streaming output), it does not optimize for scenarios where multiple inference requests are interleaved with different routing patterns. Load balancing across concurrent requests could degrade cache efficiency.

5.3.3 Cache Hit Rate Saturation

Empirical cache hit rates plateau at approximately **35%** even at relaxed capacity constraints ($C = 8$, $C = 16$). This ceiling reflects:

- **Routing Distribution Characteristics:** The expert routing distribution inherent to MoE models determines a fundamental hit rate floor. Analysis suggests that $\sim 35\%$ of expert accesses are consistently to the "hot" expert subset, while $\sim 65\%$ are to "cold" experts not resident at the time of access. This distribution is independent of cache capacity and cannot be improved through architectural changes alone.
- **Limited Temporal Locality:** While tokens exhibit some clustering (consecutive tokens tend to route to overlapping experts), the per-token routing entropy is high enough that even optimally-sized caches achieve only modest hit rates.

This hit rate ceiling means that further latency improvements beyond $C = 8$ require fundamentally different approaches (e.g., predictive prefetching or token re-ordering), not merely larger caches.

5.3.4 Staging Buffer Minimalism and Saturation Effects

The empirical observation that staging buffer count has minimal impact (within $\pm 1.2\%$ variance) suggests that:

- **Current Bottleneck:** At present, disk I/O and CPU-GPU synchronization are the primary bottlenecks, not pinned memory availability. Increasing staging buffers does not help if disk bandwidth or PCIe bandwidth remains saturated.
- **Limited Parallelism Gain:** The system achieves only partial overlap of disk I/O, PCIe transfer and computation. True pipelining would require more sophisticated scheduling, such as prefetching the next expert while the current expert is still computing (speculative loading).

5.3.5 Scope Limitations on Evaluation

The evaluation, while comprehensive, was conducted on specific hardware and models:

- **Single Hardware Generation:** Experiments used NVIDIA T4 and H100 GPUs. Results on AMD ROCm devices, Intel Arc, or older NVIDIA hardware (V100, A100) may differ significantly due to different PCIe bandwidth, memory subsystem characteristics and driver optimization.
- **Limited Model Coverage:** Primary evaluation focused on Granite-3B. While theoretical analysis covers Qwen2-MoE and JetMoE, empirical validation on these larger models was not performed.
- **Artificial Workloads:** The vLLM PagedAttention benchmark, while standard, may not reflect all production inference patterns (e.g., long-context requests, bursty token arrival patterns, or adversarial expert routing distributions).

5.4 Future Work

The foundation established by CacheFlow opens several promising research directions:

5.4.1 Predictive Prefetching Based on Token Embeddings

Current prefetching is reactive—experts are fetched only after a cache miss is detected. A predictive approach would:

- **Analyze Token Embeddings:** Compute semantic similarity between the current token and upcoming tokens to predict which experts are likely to be requested in the next few time steps.
- **Speculative Loading:** Initiate disk reads for predicted experts while current experts are still being computed, reducing or eliminating the cache miss penalty for predicted experts.
- **Adaptive Thresholds:** Adjust prefetching aggressiveness based on prediction confidence and available disk bandwidth.

This approach could push cache hit rates from $\sim 35\%$ toward 50–60% for tokens where semantic clustering is strong, with minimal false-positive penalty for mispredictions.

5.4.2 Dynamic and Adaptive Staging Buffer Sizing

Current staging buffer allocation is static. Dynamic sizing would:

- **Monitor I/O Characteristics:** Track disk I/O latency, PCIe transfer throughput and GPU computation time to dynamically adjust `staging_slots` based on current bottleneck.
- **Reduce Memory Overhead:** When disk bandwidth is the bottleneck (as is currently the case), reduce staging buffers to free pinned memory for other uses. When GPU computation is the bottleneck, increase staging buffers to enable more aggressive prefetching.
- **Online Adaptation:** Implement a feedback loop that adjusts staging buffer count without restarting the inference engine, enabling real-time optimization based on workload characteristics.

5.4.3 Multi-GPU Cache Coherency and Expert Parallelism

Extending CacheFlow to multi-GPU setups requires:

- **Distributed Cache State:** Implement a coherency protocol (e.g., invalidation-based or write-through) to prevent redundant expert fetches across GPUs. When Expert 5 is loaded on GPU 0, GPU 1 should be aware and potentially request it via GPU-to-GPU direct memory access (NVLink or PCIe) rather than reloading from disk.
- **Expert Parallelism Scheduling:** For models with multiple disjoint expert groups, partition experts across GPUs and implement cross-GPU communication for token routing and expert output aggregation.
- **Load Balancing:** Design scheduling algorithms that account for heterogeneous expert assignment across GPUs and load-imbalance due to routing variation.

This extension would enable scaling CacheFlow to 8–16 GPU clusters, potentially enabling even larger MoE models to be deployed efficiently.

5.4.4 Token Reordering and Routing-Aware Batching

Current batching processes tokens in arrival order. Reordering could:

- **Expert-Aligned Batching:** Reorder tokens within a batch such that tokens requiring the same experts are processed consecutively, maximizing cache hits and minimizing context switches.
- **Deadlock Avoidance:** Design reordering algorithms that prevent expert scheduling deadlocks (e.g., when all cached experts are frozen and all cold experts require eviction).

- **Latency Trade-offs:** Balance latency improvements from reordering against increased queuing delay for tokens waiting for reordering to complete.

5.4.5 Hardware-Specific Optimizations

- **Direct Disk-to-GPU Transfers:** Explore GPU direct-attached NVMe (e.g., via NVLink or PCIe peer-to-peer) to bypass host CPU, reducing latency from ~ 150 ms to ~ 50 ms per expert.
- **Lossless Compression:** Implement on-the-fly decompression of expert weights (e.g., using block-floating-point or low-rank factorization) to reduce disk I/O volume and transfer time by 20–40%.
- **Custom Kernels:** Develop GPU kernels for expert loading and cache management, reducing CPU-GPU synchronization overhead.

5.4.6 Theoretical Analysis of Cache Optimality

- **Hit Rate Bounds:** Derive theoretical bounds on achievable cache hit rates as functions of routing distribution entropy and capacity constraint, enabling principled prediction of system performance.
- **Scheduling Complexity:** Analyze the computational complexity of optimal prefetching and eviction scheduling, potentially uncovering approximation algorithms with better performance guarantees than LRU.

5.5 Concluding Remarks

This thesis has demonstrated that intelligent systems-level optimization can transform Mixture-of-Experts models from theoretical constructs accessible only to well-funded research institutions into practical, deployable systems for mainstream researchers and practitioners. CacheFlow’s three-tier memory hierarchy—seamlessly chaining disk-resident experts through pinned CPU staging to GPU VRAM slots—proves that GPU memory constraints need not preclude deployment of state-of-the-art sparse architectures.

The empirical achievements are substantial: a 5–9 $\times$ reduction in GPU memory footprint while maintaining reasonable latencies (600 ms at ultra-constrained settings, improving to 460 ms at relaxed constraints) and throughputs approaching 130 tokens/minute that enable practical inference workloads. More importantly, these achievements unlock access to 60+ expert models on consumer-grade hardware (12 GB VRAM) that previously required specialized infrastructure (80+ GB VRAM).

Beyond the specific technical contributions, this work addresses a critical democratization gap in the AI landscape. Large language models trained with MoE architectures represent the frontier of AI research, yet their deployment has been restricted to organizations with sufficient computational budget. By enabling MoE inference on modest hardware without sacrificing computational correctness or absolute performance, CacheFlow expands the constituency of researchers and developers who can experiment with, deploy and advance these powerful models.

The limitations identified—I/O bandwidth bottlenecks, single-GPU scope and routing distribution constraints—are not mere technical shortcomings but natural boundaries of the approach. They point toward future research directions: predictive prefetching, dynamic buffer sizing, multi-GPU coherency and hardware-specific optimizations that can further close the gap between constrained and unlimited-memory inference.

Ultimately, CacheFlow represents a fundamental shift in how sparse neural architectures can be deployed. By decoupling GPU memory requirements from model size through intelligent scheduling and caching, the system opens new possibilities for scalable, efficient and accessible AI inference. As MoE models continue to proliferate in both research and production settings, systems-level optimizations like CacheFlow will be essential infrastructure for ensuring that the benefits of these architectures can be widely realized, not merely in academic publications but in practical deployments across the global research and practitioner community.

References

- [1] R. A. Jacobs, M. I. Jordan, S. J. Nowlan, and G. E. Hinton, “Adaptive mixtures of local experts,” *Neural Computation*, 1991.
- [2] M. I. Jordan and R. A. Jacobs, “Hierarchical mixtures of experts and the em algorithm,” *Neural Computation*, 1994.
- [3] D. P. Kingma and J. Ba, “Adam: A method for stochastic optimization,” *arXiv preprint arXiv:1412.6980*, 2014.
- [4] G. Hinton, O. Vinyals, and J. Dean, “Distilling the knowledge in a neural network,” *arXiv preprint arXiv:1503.02531*, 2015.
- [5] D. Bahdanau, K. Cho, and Y. Bengio, “Neural machine translation by jointly learning to align and translate,” in *International Conference on Learning Representations (ICLR)*, 2015.
- [6] S. Ioffe and C. Szegedy, “Batch normalization: Accelerating deep network training by reducing internal covariate shift,” in *International Conference on Machine Learning (ICML)*, 2015.
- [7] L. J. Ba, J. R. Kiros, and G. E. Hinton, “Layer normalization,” *arXiv preprint arXiv:1607.06450*, 2016.
- [8] D. Hendrycks and K. Gimpel, “Gaussian error linear units (gelus),” *arXiv preprint arXiv:1606.08415*, 2016.
- [9] A. Vaswani, N. Shazeer, N. Parmar, *et al.*, “Attention is all you need,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.
- [10] N. Shazeer, A. Mirhoseini, K. Maziarz, *et al.*, “Outrageously large neural networks: The sparsely-gated mixture-of-experts layer,” in *International Conference on Learning Representations (ICLR)*, 2017.
- [11] E. Grave, A. Joulin, M. Cissé, D. Grangier, and H. Jégou, “Efficient softmax approximation for gpus,” in *International Conference on Machine Learning (ICML)*, 2017.

- [12] A. Radford, K. Narasimhan, T. Salimans, and I. Sutskever, “Improving language understanding by generative pre-training,” tech. rep., OpenAI, 2018.
- [13] Z. Yang, Z. Dai, R. Salakhutdinov, and W. W. Cohen, “Breaking the softmax bottleneck: A high-rank rnn language model,” in *International Conference on Learning Representations (ICLR)*, 2018.
- [14] A. Radford, J. Wu, R. Child, *et al.*, “Language models are unsupervised multitask learners,” tech. rep., OpenAI, 2019.
- [15] R. Child, S. Gray, A. Radford, and I. Sutskever, “Generating long sequences with sparse transformers,” in *arXiv preprint arXiv:1904.10509*, 2019.
- [16] B. Zhang and R. Sennrich, “Root mean square layer normalization,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.
- [17] T. B. Brown, B. Mann, N. Ryder, *et al.*, “Language models are few-shot learners,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
- [18] C. Raffel, N. Shazeer, A. Roberts, *et al.*, “Exploring the limits of transfer learning with a unified text-to-text transformer,” *Journal of Machine Learning Research (JMLR)*, 2020.
- [19] J. Rasley, S. Rajbhandari, O. Ruwase, and Y. He, “Deepspeed: System optimizations enable training deep learning models with over 100 billion parameters,” in *Proceedings of the 26th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining (KDD)*, 2020.
- [20] A. Dosovitskiy *et al.*, “An image is worth 16x16 words: Transformers for image recognition at scale,” in *International Conference on Learning Representations (ICLR)*, 2020.
- [21] N. Shazeer, “Glu variants improve transformer,” *arXiv preprint arXiv:2002.05202*, 2020.
- [22] I. Beltagy, M. E. Peters, and A. Cohan, “Longformer: The long-document transformer,” *arXiv preprint arXiv:2004.05150*, 2020.
- [23] Z. Lan, M. Chen, S. Goodman, K. Gimpel, P. Sharma, and R. Soricut, “Albert: A lite bert for self-supervised learning of language representations,” in *International Conference on Learning Representations (ICLR)*, 2020.
- [24] D. Lepikhin, H. Lee, Y. Xu, *et al.*, “Gshard: Scaling giant models with conditional computation and automatic sharding,” in *International Conference on Learning Representations (ICLR)*, 2021.

- [25] W. Fedus, B. Zoph, and N. Shazeer, “Switch transformers: Scaling to trillion parameter models with simple and efficient sparsity,” *Journal of Machine Learning Research (JMLR)*, 2022.
- [26] T. Dao, D. Y. Fu, S. Ermon, A. Rudra, and C. Ré, “Flashattention: Fast and memory-efficient exact attention with io-awareness,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
- [27] W. Kwon, Z. Li, S. Zhuang, Y. Sheng, L. Zheng, C. H. Yu, J. E. Gonzalez, H. Zhang, and I. Stoica, “Efficient memory management for large language model serving with pagedattention,” in *ACM Symposium on Operating Systems Principles (SOSP)*, 2023.
- [28] H. Touvron, T. Lavril, G. Izacard, *et al.*, “Llama: Open and efficient foundation language models,” *arXiv preprint arXiv:2302.13971*, 2023.
- [29] H. Yoo and H. Lee, “Swapmoe: Serving off-the-shelf moe-based large language models with tunable memory budget,” *arXiv preprint arXiv:2309.09515*, 2023.
- [30] S. Zuo *et al.*, “Moe-infinity: Efficient moe inference on personal machines with sparsity-aware expert cache,” *arXiv preprint arXiv:2305.11794*, 2023.
- [31] Microsoft, “Deepspeed-mii: Model inference interface,” 2023.
- [32] NVIDIA, “Tensorrt-llm: Optimized inference for large language models,” 2023.
- [33] T. Dao, “Flashattention-2: Faster attention with better parallelism and work partitioning,” *arXiv preprint arXiv:2307.08691*, 2023.
- [34] W. Liang *et al.*, “Deepseek-v2: A strong, economical, and efficient mixture-of-experts language model,” *arXiv preprint arXiv:2405.04434*, 2024.
- [35] X. Shen *et al.*, “fmoe: Fine-grained and prompt-aware expert offloading for large mixture-of-experts serving,” *arXiv preprint arXiv:2403.12881*, 2024.